

BE 537 - Assignment 2

November 13, 2022

1 Introduction

We will be implementing the non-symmetric version of the *Log-Domain Demons* algorithm by Vercauteren et al. (2008). Before the advent of deep learning image registration, the approach was arguably the most widely used deformable image registration approach because of its speed and ability to produce diffeomorphic transformations. You may find it easiest to follow the description of the algorithm in Section 3.2.3 (Algorithm 4) in the [Non-Parametric Volumetric Registration \(NVR\) book chapter](#).

In this assignment, most of the code is already provided. There are six tasks in the base component of the assignment that require answering mathematical questions or coding, indicated in bold below. The total number of points for these seven tasks is 70.

2 Preliminaries

```
[1]: from google.colab import drive
drive.mount('/content/drive')

import sys
sys.path.insert(0, 'drive/MyDrive/be537/assignments/assignment2/package/')
```

Mounted at /content/drive

2.1 Load required packages

I have provided an implementation of the functions from the first assignment in a Python module `be537hw1`. You will use some of these functions again in this assignment. There are also a few helpful functions in module `be537hw2` that we will use. To see the listing of functions available in `be537hw1` and `be537hw2`, type

```
help('be537hw1')
help('be537hw2')
```

The cell below loads the required packages and sets the data path. **Update the data path for your environment.**

```
[2]: # Import required libraries
import os
import numpy as np
```

```

import nibabel as nib
import matplotlib.pyplot as plt
import torch
import torch.nn.functional as tfun

# Import the functions from the first assignment
from be537hw1 import *
from be537hw2 import *

# Configure matplotlib options
import matplotlib
%matplotlib inline
matplotlib.rcParams['figure.figsize'] = [10, 9]

%precision %.4f

# Path to the data for this experiment
data_path='drive/MyDrive/be537/assignments/assignment2/package/data/'

```

2.2 Set up GPU capabilities

The cell below sets up a CUDA device to use with torch, if available to you. If you want to use the CPU even if a GPU is present, set `force_cpu` to be `True`.

You can create tensors that reside on the GPU with the command `T = torch.Tensor(..., device=cuda)` and copy existing tensors to the GPU with the command `T.to(cuda)`. If you do not have a GPU, these commands will still work.

- You do not need a GPU for the base component of the assignment, but I recommend using it for the extensions

```

[3]: # Set up CUDA if available
force_cpu = False
if torch.cuda.is_available() and not force_cpu:
    cuda = torch.device('cuda', 0)
    print("Using CUDA device %d named '%s' with %6.2f Gb total memory" %
          (cuda.index, torch.cuda.get_device_name(cuda.index),
           torch.cuda.get_device_properties(cuda.index).total_memory / 2.0**30))
else:
    cuda = torch.device('cpu')
    print("Using CPU")

```

Using CUDA device 0 named 'Tesla T4' with 14.76 Gb total memory

2.3 New Visualization Code

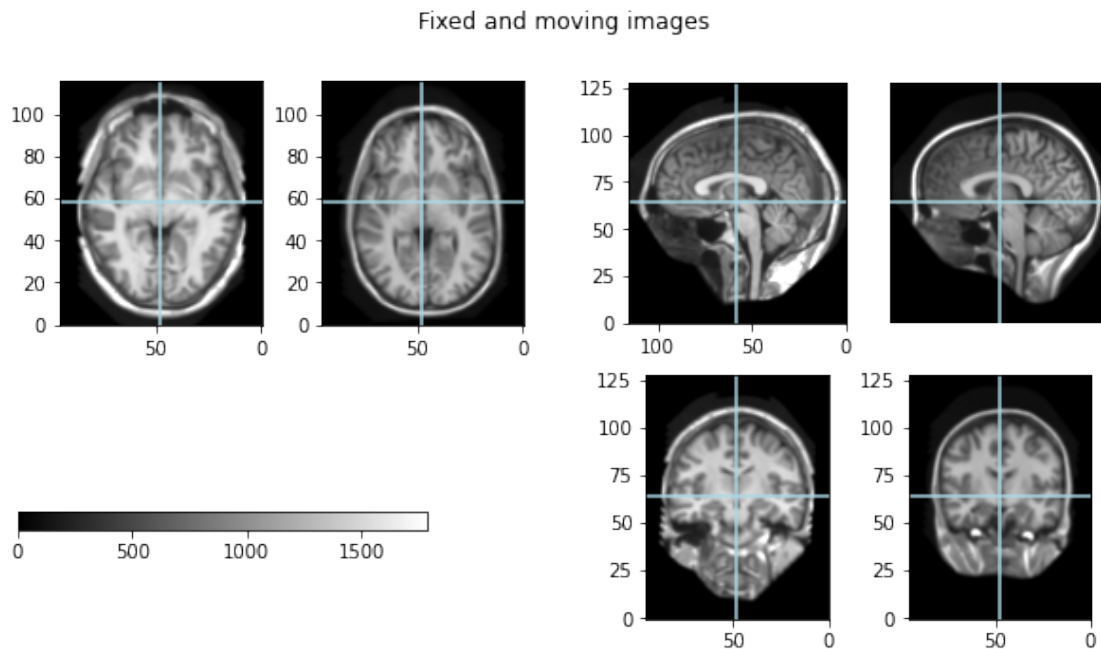
Module `be537hw2` contains two new functions `my_view_multi` and `my_view_ovl` for visualizing multiple images at once. They will be useful for viewing registration results. The code below loads

a pair of images from the `testing` directory. The first is the atlas image with ID *1000*, which we will be treating as the *fixed* image. The second is the atlas with ID *1001* that has been aligned to the fixed image using affine registration. We will be treating it as the *moving* image for our deformable registration experiments.

```
[4]: # Read the fixed image
I_fix, hdr_fix = my_read_pytorch_image_from_nifti(
    os.path.join(data_path, 'testing/atlas_2mm_1000_3.nii.gz'),
    dtype=torch.float32, device=cuda)


# Read the moving image
I_mov, hdr_mov = my_read_pytorch_image_from_nifti(
    os.path.join(data_path, 'testing/reslice_affine_fx_1000_mv_1001.nii.gz'),
    dtype=torch.float32, device=cuda)

# View the fixed and moving images side by side
my_view_multi([I_fix, I_mov], hdr_fix, cmap='gray', title='Fixed and moving_
↳images')
```



And here the two images are displayed as overlays

```
[5]: my_view_overlay(I_fix, I_mov, hdr_fix, cmap='gray')
```



2.4 Working with deformable transformations

Before getting started, it is important to understand how spatial transformations are represented in Pytorch.

2.4.1 Order of dimensions

Recall that 3D images are represented in Pytorch as 5D-tensors of dimensions $[N, C, D, H, W]$ where N is the mini-batch size (1 for now), C is the number of channels (1 for grayscale, 3 for an RGB image), and D, H, W are the dimensions of the 3D image.

Spatial transformations are 3D images that contain a coordinate $\phi(x)$ at every voxel x . They can be represented as 3D images with $C = 3$ channels. However, recall from Assignment 1 that the function `grid_sample` expects spatial transformations to be represented as 5D tensors with dimensions $[N, D, H, W, 3]$. Notice that instead of storing the spatial coordinates in the second

dimension, the coordinates are stored in the last dimension. In this assignment, we will represent spatial transformations consistent with `grid_sample`.

- You can use function `torch.permute` to go back and forth between $[N, 3, D, H, W]$ and $[N, D, H, W, 3]$ representations of a spatial transformation

2.4.2 Spatial Transformation vs. Displacement Field

Recall the distinction between a *spatial transformation* and a *displacement field*. A transformation ϕ assigns to each point x in the image domain another point, $\phi(x)$. A displacement field u assigns to each point x a displacement vector $u(x)$. The two are related simply as $\phi(x) = x + u(x)$. Because of this simple relationship, we can think of a displacement field as just another way to represent a spatial transformation. It turns out that in practice representing spatial transformations as displacement fields is easier, particularly when it comes to applying composition operations to spatial transformations. In this assignment, we will represent spatial transformations as displacement fields.

2.4.3 Displacement Field IO

Module `be537hw2` contains the function `my_read_pytorch_warp_from_nifti` that can read a displacement field output by registration tools [ANTs](#) and [Greedy](#) and return a PyTorch tensor compatible with `grid_sample`.


The next code snippet loads a displacement field (or “warp”) computed using [Greedy](#) between the fixed and moving images and applies it to the moving image. The transformed (or “resliced”) moving image is displayed as an overlay on top of the fixed image. Observe that the contours from the resliced moving image match the structures in the fixed image better than in the previous plot.

```
[6]: # Read the spatial transformation (warp for short)
I_warp, hdr_warp = my_read_pytorch_warp_from_nifti(
    os.path.join(data_path, 'testing/warp_fx_1000_mv_1001.nii.gz'),
    dtype=torch.float32, device=cuda)

# Create a coordinate grid for using with grid_sample
grid = my_create_pytorch_grid(I_fix.shape).to(cuda)

# Reslice the moving image using the warp
I_res = tfun.grid_sample(I_mov, grid + I_warp, align_corners=False)

# View the fixed and resliced images
my_view_overlay(I_fix, I_res, hdr_fix, cmap='gray')
```



3 Implement Diffeomorphic Demons Registration in PyTorch

In this section, we will implement the diffeomorphic demons image registration algorithm and test it on the imaging data provided. Most code is provided. The portions where you have to write code or answer questions are labeled **TASK 1**, etc.

3.1 Building blocks of the algorithm

Implementing the *Log-Domain Demons* algorithm by Vercauteren et al. (2008) will require the following components:

- **Scaling and squaring**, a technique that converts a smooth and bounded velocity field into a diffeomorphic transformation
- **Image gradient computation**, needed to compute optical flow between fixed and moving

images

- **Gaussian smoothing**, used to regularize velocity fields and displacement fields
- **Optical flow computation**,

3.2 Scaling and squaring

In this part of the assignment, you will implement the scaling and squaring operation. The scaling and squaring algorithm is described in Section 2.1.4 of the NVR chapter. We also discuss it in class on the first day of the deformable registration lecture. Scaling and squaring takes as input a vector field called a **stationary velocity field (SVF)** and generates a spatial transformation that is diffeomorphic.

3.2.1 Theory: Scaling and Squaring using Displacement Fields

The scaling and squaring algorithm is described in Section 2.1.4 of the NVR book chapter in terms of compositions of spatial transformations ψ_k . Since we will be working with displacement fields, we need to rewrite the algorithm in terms of displacement fields.

In a Markdown cell in your notebook, prove the following:

- If ψ_1 and ψ_2 are transformations with corresponding displacement fields v_1 and v_2 , then the displacement field w corresponding to the transformation $\psi_1 \circ \psi_2$ (the composition of ψ_1 and ψ_2) is given by

$$w(x) = v_2(x) + v_1(x + v_2(x))$$

- Next, rewrite the scaling and squaring algorithm (at the end of section 2.1.4 of the NVR chapter) using displacement fields v instead of spatial transformations ψ .

TASK 1a (5 pts): Your proof goes here

$$\psi_1(x) = x + v_1(x)$$

$$\psi_2(x) = x + v_2(x)$$

Let,

$$\psi_3(x) = \psi_1(x) \circ \psi_2(x) = \psi_1(\psi_2(x))$$

$$\psi_3(x) = \psi_1(x + v_2(x))$$

$$\therefore \psi_3(x) = x + v_2(x) + v_1(x + v_2(x))$$

The displacement field of $\psi_3(x)$ is $\psi_3(x) - x$

$$\therefore w(x) = v_2(x) + v_1(x + v_2(x))$$

TASK 1b (5 pts): Your algorithm goes here

Let, $u(x)$ be a stationary velocity field.

With n iterations,

$$\psi_0(x) = x + \frac{1}{2^n} u(x)$$

$$v_0(x) + x = x + \frac{1}{2^n} u(x)$$

So initialization:

$$v_0(x) = \frac{1}{2^n} u(x)$$

The algorithm for k-th iteration:

$$v_k = v_{k-1}(x) + v_{k-1}(x + v_{k-1}(x)) \text{ for } k = 1, \dots, K$$

3.2.2 Implementation of Scaling and Squaring

TASK 2 (10 pts): Complete the code for a function that performs scaling and squaring.

Hints: - The only PyTorch functions needed to implement this function are `grid_sample` and `permute`. - When calling `grid_sample`, the first input will need to be of shape `[1,3,D,H,W]` and the output will be of shape `[1,3,D,H,W]`. Since `grid_sample` expects inputs to have shape `[1,D,H,W,3]`, you will need to use `permute` twice inside of this function.

```
[7]: def my_scaling_and_squaring(u, grid, n = 6):
    """
    Apply scaling and squaring to a displacement field

    The input and output to this method are displacement fields, v and u,
    respectively. If u is sufficiently smooth and bounded, theoretically,
    the deformation field  $\phi(x) = x + u(x)$  will be diffeomorphic

    :param u: Input stationary velocity field, PyTorch tensor of shape
    ↪ [1,D,H,W,3]
    :param grid: Sampling grid of size [1,D,H,W,3] (created by
    ↪ my_create_pytorch_grid)
    :param n: Number of iterations of scaling and squaring (default: 6)

    :returns: Output displacement field, v, PyTorch tensor of shape [1,D,H,W,3]
    """
    v0 = u*((0.5)**n) # initialize
    vi = torch.permute(v0, (0,4,1,2,3))

    for _ in range(n):
        curr_x = grid + torch.permute(vi, (0,2,3,4,1))

        # recursion
        vi = vi + tfun.grid_sample(vi, curr_x, align_corners=False)

    v = torch.permute(vi, (0,2,3,4,1))

    return v
```

3.2.3 Test Scaling and Squaring

To confirm that your code is working correctly, we perform the following test:

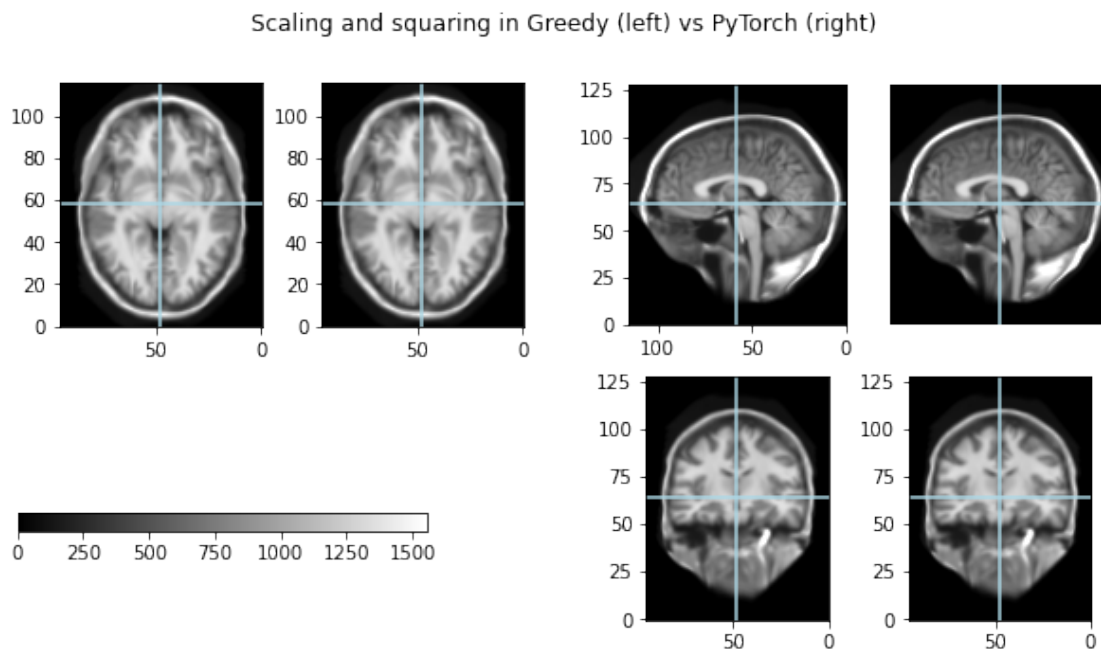
1. Load a sample stationary velocity field `I_svf` from `testing/svf_fx_1000_mv_1001.nii.gz`
2. Apply scaling and squaring to `I_svf` and use the result to transform the moving image
3. Compare the result to the moving image transformed using `I_warp` above

```
[8]: # Load the root warp
I_svf, hdr_svf = my_read_pytorch_warp_from_nifti(os.path.
    ↪join(data_path, 'testing/svf_fx_1000_mv_1001.nii.gz'), dtype=torch.float32, ↪
    ↪device=cuda)

# Apply scaling and squaring to the SVF
I_warp_ss = my_scaling_and_squaring(I_svf, grid, 6)

# Sample the moving image using this transformation
I_res_ss = tfun.grid_sample(I_mov, grid + I_warp_ss, align_corners=False)

# Compare the two images
my_view_multi([I_res, I_res_ss], hdr_fix, cmap='gray',
    title='Scaling and squaring in Greedy (left) vs PyTorch (right)')
```



Here is the quantitative comparison of the two images (you should expect the values to be ~ 0.1 for mean absolute difference and ~ 4.0 for max absolute difference)

```
[9]: print('Difference between I_res and I_res_ss')
print(' Mean absolute difference: ', torch.mean(torch.abs(I_res-I_res_ss)))
print(' Maximum absolute difference: ', torch.max(torch.abs(I_res - I_res_ss)))
```

Difference between `I_res` and `I_res_ss`

Mean absolute difference: tensor(0.1036, device='cuda:0')
Maximum absolute difference: tensor(4.2585, device='cuda:0')

3.3 Gradient computation

We will need to approximate the gradient of the moving image during each step of the deformable registration algorithm. The gradient of an image can be represented as a multi-component image $\nabla I(x) : \mathbb{R}^3 \rightarrow \mathbb{R}^3$. At each voxel, this image contains the three partial derivatives of I with respect to the coordinates x_1, x_2, x_3 :

$$\nabla I(x) = \left[\frac{\partial I}{\partial x_1}, \frac{\partial I}{\partial x_2}, \frac{\partial I}{\partial x_3} \right]^T$$

It turns out that the images $\frac{\partial I}{\partial x_1}$, $\frac{\partial I}{\partial x_2}$, and $\frac{\partial I}{\partial x_3}$ can each be approximated as discrete convolutions of the image I with simple kernels consisting of three elements.

TASK 3 (10 pts): Your derivation of the convolution kernels for image gradient computation goes here

$$I(y) \simeq I(x) + \sum_{i=1}^3 \frac{\partial I}{\partial x_i}(x) \cdot (y_i - x_i) + \sum_{i=1}^3 \sum_{j=1}^3 \frac{\partial^2 I}{\partial x_i \partial x_j}(x) \cdot \frac{(y_i - x_i)(y_j - x_j)}{2}$$

For $\frac{\partial I}{\partial x_1}$, since we are interested in the $i = 1$ term of x :

$$I(x + \delta_1 e_1) = I(x) + \frac{\partial I}{\partial x_1}(x)(x_1 + \delta_1 - x_1) + \frac{\partial^2 I}{\partial x_1^2}(x) \frac{(x_1 + \delta_1 - x_1)(x_1 + \delta_1 - x_1)}{2}$$

$$\therefore I(x + \delta_1 e_1) = I(x) + \delta_1 \frac{\partial I}{\partial x_1}(x) + \frac{\delta_1^2}{2} \frac{\partial^2 I}{\partial x_1^2}(x)$$

And,

$$I(x - \delta_1 e_1) = I(x) + \frac{\partial I}{\partial x_1}(x)(x_1 - \delta_1 - x_1) + \frac{\partial^2 I}{\partial x_1^2}(x) \frac{(x_1 - \delta_1 - x_1)(x_1 - \delta_1 - x_1)}{2}$$

$$\therefore I(x - \delta_1 e_1) = I(x) - \delta_1 \frac{\partial I}{\partial x_1}(x) + \frac{\delta_1^2}{2} \frac{\partial^2 I}{\partial x_1^2}(x)$$

So,

$$I(x + \delta_1 e_1) - I(x - \delta_1 e_1) = I(x) + \delta_1 \frac{\partial I}{\partial x_1}(x) + \frac{\delta_1^2}{2} \frac{\partial^2 I}{\partial x_1^2}(x) - I(x) + \delta_1 \frac{\partial I}{\partial x_1}(x) - \frac{\delta_1^2}{2} \frac{\partial^2 I}{\partial x_1^2}(x)$$

$$I(x + \delta_1 e_1) - I(x - \delta_1 e_1) = 2\delta_1 \frac{\partial I}{\partial x_1}(x)$$

$$\therefore \frac{\partial I}{\partial x_1}(x) = -\frac{1}{2\delta_1} I(x - \delta_1 e_1) + \frac{1}{2\delta_1} I(x + \delta_1 e_1)$$

If we consider $d_1 = -\frac{1}{2\delta_1}$ and $d_2 = \frac{1}{2\delta_1}$, this relationship between the derivative of $I(x)$ at each pixel and the value of $I(x)$ at adjacent pixels can also be expressed as a discrete convolution between $I(x)$ and a size (3,1,1) kernel $[[[d_1]], [[0]], [[d_2]]]$.

Similarly for $\frac{\partial I}{\partial x_2}$, we have:

$$\therefore \frac{\partial I}{\partial x_2}(x) = -\frac{1}{2\delta_2} I(x - \delta_2 e_2) + \frac{1}{2\delta_2} I(x + \delta_2 e_2)$$

$$h_1 = -\frac{1}{2\delta_2} \text{ and } h_2 = \frac{1}{2\delta_2}$$

The relationship between the derivative of $I(x)$ at each pixel and the value of $I(x)$ at adjacent pixels can also be expressed as a discrete convolution between $I(x)$ and a size (1,3,1) kernel $[[[h_1]], [0], [[h_2]]]$

For $\frac{\partial I}{\partial x_3}$, we have:

$$\therefore \frac{\partial I}{\partial x_3}(x) = -\frac{1}{2\delta_3}I(x - \delta_3 e_3) + \frac{1}{2\delta_3}I(x + \delta_3 e_3)$$

$$w_1 = -\frac{1}{2\delta_3} \text{ and } w_2 = \frac{1}{2\delta_3}$$

The relationship between the derivative of $I(x)$ at each pixel and the value of $I(x)$ at adjacent pixels can also be expressed as a discrete convolution between $I(x)$ and a size (1,1,3) kernel $[[[w_1, 0, w_2]]]$

3.3.1 Implementation

TASK 4 (10 pts): Write a function with a signature below that computes the gradient of a 3D image

Hints: * The gradient computation can be accomplished using a single call to `torch.nn.functional.conv3d` with `I` as the first parameter and the second parameter a tensor of size `[3, 1, 3, 3, 3]` that contains your three kernels. * The spacings $\delta_1, \delta_2, \delta_3$ should be with respect to the PyTorch coordinate system, in which the image domain is the cube $[-1, 1] \times [-1, 1] \times [-1, 1]$. So for an image of size $N_1 \times N_2 \times N_3$, we will have $\delta_1 = \frac{2}{N_1}$, etc.

```
[10]: # Compute the gradient of a 3D image in PyTorch
def my_image_gradient(I):
    """
    Compute the gradient of an image using central difference approximation

    :param I: input image, represented as a [N,1,D,H,W] tensor
    :returns: gradient of the input image, represented as a [N,3,D,H,W] tensor
    """
    (N, C, D, H, W) = I.shape

    k1 = torch.tensor([-0.25*D, 0, 0.25*D], device=cuda).reshape((3,1,1))
    k2 = torch.tensor([-0.25*H, 0, 0.25*H], device=cuda).reshape((1,3,1))
    k3 = torch.tensor([-0.25*W, 0, 0.25*W], device=cuda).reshape((1,1,3))

    k1 = tfun.pad(k1, (1, 1, 1, 1, 0, 0))
    k2 = tfun.pad(k2, (1, 1, 0, 0, 1, 1))
    k3 = tfun.pad(k3, (0, 0, 1, 1, 1, 1))

    kernel = torch.stack((k3, k2, k1)).unsqueeze(0).reshape((3,1,3,3,3))

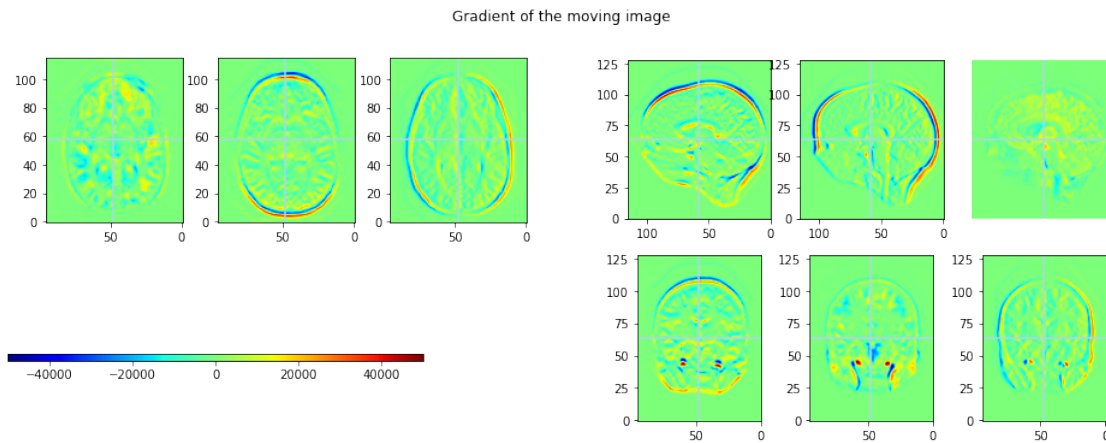
    I_grad = tfun.conv3d(I, kernel, padding="same")

    return I_grad
```

3.3.2 Check Gradient Computation

Plot the image gradient using the commands. The expected output is shown in files/gradient_example.png.


```
[11]: I_mov_grad = my_image_gradient(I_mov)
my_view_multi(I_mov_grad, hdr_mov, cmap='jet', crange=[-50000,50000],
              transpose=True, title='Gradient of the moving image',
              figwidth=16);
```



Also plot the gradient magnitude image i.e. an image where at every voxel the intensity is equal to $|\nabla I|$. The gradient magnitude is large at edges in the image.

```
[12]: my_view_multi(torch.sqrt(I_mov_grad.square().sum(1)), hdr_mov, cmap='jet',
                    title='Gradient magnitude of the moving image', figwidth=6)
```

Gradient magnitude of the moving image



3.4 Gaussian Smoothing

Implementing Gaussian smoothing in Pytorch efficiently requires computing the fast Fourier transform (FFT) of the input image and the Gaussian kernel, multiplication in the frequency domain, and taking the inverse FFT. Getting this to work correctly is a little tricky, so the code for Gaussian smoothing is provided for you in functions `my_pytorch_gaussian_lpf` and `my_pytorch_gaussian_lpf_transform`.

```
[13]: # Function to generate FFT of a Gaussian kernel for use in PyTorch FFT
def my_gaussian_fft_kernel(I_ref, sigma):
    """
    Generate the FFT of a Gaussian kernel that can be used for FFT-based
    ↪ Gaussian smoothing
```

```

:param I_ref: 5D tensor that will be convolved with the kernel, used to
↳infer the
           shape, data type and device for the output kernel
:param sigma: standard deviation of the Gaussian kernel
:param dtype: data type of the output tensor (default: torch.float32)
:returns: 5D tensor that can be multiplied by the image RFFT for FFT-based
↳smoothing
"""

# Generate a Gaussian kernel, convert to PyTorch
K = torch.tensor(my_gaussian_3d(sigma),
                 dtype=I_ref.dtype,
                 device=I_ref.device).unsqueeze(0).unsqueeze(0)

# Normalize the kernel
K = K / K.sum()

# Pad the kernel so it is centered in an image of size I_ref.shape
ab_pad = np.array(I_ref.shape[2:5]) - np.array(K.shape[2:5])
a_pad = np.floor(ab_pad / 2)
b_pad = ab_pad - a_pad
pad_5d = tuple(np.flip(np.vstack((a_pad, b_pad)).T.flatten()).
↳astype(int))+(0,0,0,0)
K_pad = tfun.pad(K,pad_5d)

# Return the real FFT of the kernel
return torch.fft.rfftn(K_pad)

# Function to smooth a scalar image with a Gaussian using PyTorch FFT
def my_pytorch_gaussian_lpf(img, sigma=None, kernel=None):
    """
    Apply Gaussian smoothing to a 3D image represented as a PyTorch tensor

    You can pass in the sigma of the Gaussian or a 5D tensor representing the
    ↳Gaussian
    kernel (generated using `my_gaussian_fft_kernel`). The latter is faster if
    ↳you will
    be making repeated calls to this function.

    :param img: Input image, represented as a 5D tensor
    :param sigma: Standard deviation of the Gaussian kernel. If `kernel`
    ↳parameter
    is not supplied, the kernel will be generated by calling
    ↳`my_gaussian_fft_kernel`

```

```

:param kernel: Gaussian kernel generated by `my_gaussian_fft_kernel`.
"""

# Either sigma or kernel must be provided
assert sigma is not None or kernel is not None
if kernel is None:
    kernel = my_gaussian_fft_kernel(img, sigma)

# Create a Gaussian kernel
img_fft = torch.fft.rfftn(img)
return torch.fft.fftshift(torch.fft.irfftn(img_fft * kernel))

# Function to smooth a deformation field with a Gaussian
def my_pytorch_gaussian_lpf_transform(v, sigma=None, kernel=None):
    """
    Apply Gaussian smoothing to a transform represented as a PyTorch tensor

    You can pass in the sigma of the Gaussian or a 5D tensor representing the
    ↪ Gaussian
    kernel (generated using `my_gaussian_fft_kernel`). The latter is faster if
    ↪ you will
    be making repeated calls to this function.

    :param img: Input spatial transformation, represented as a [N,D,H,W,3]
    ↪ tensor
    :param sigma: Standard deviation of the Gaussian kernel. If `kernel`
    ↪ parameter
    is not supplied, the kernel will be generated by calling
    ↪ `my_gaussian_fft_kernel`
    :param kernel: Gaussian kernel generated by `my_gaussian_fft_kernel`.
    """

    # Make sure the image is formatted as a transform
    assert v.shape[4]==3

    # Convert the transform into a multi-channel image of shape [1,3,D,H,W]
    v_chan = v.permute(0,4,1,2,3)

    # Either sigma or kernel must be provided
    assert sigma is not None or kernel is not None
    if kernel is None:
        kernel = my_gaussian_fft_kernel(v_chan, sigma)

    # Perform FFT convolution but only along image dimensions
    v_fft = torch.fft.rfftn(v_chan, dim=(-3,-2,-1))
    v_sm_chan = torch.fft.fftshift(

```

```

torch.fft.irfftn(v_fft * kernel, dim=(-3,-2,-1)),
dim=(-3,-2,-1))

# Reorder dimensions again
return v_sm_chan.permute(0,2,3,4,1)

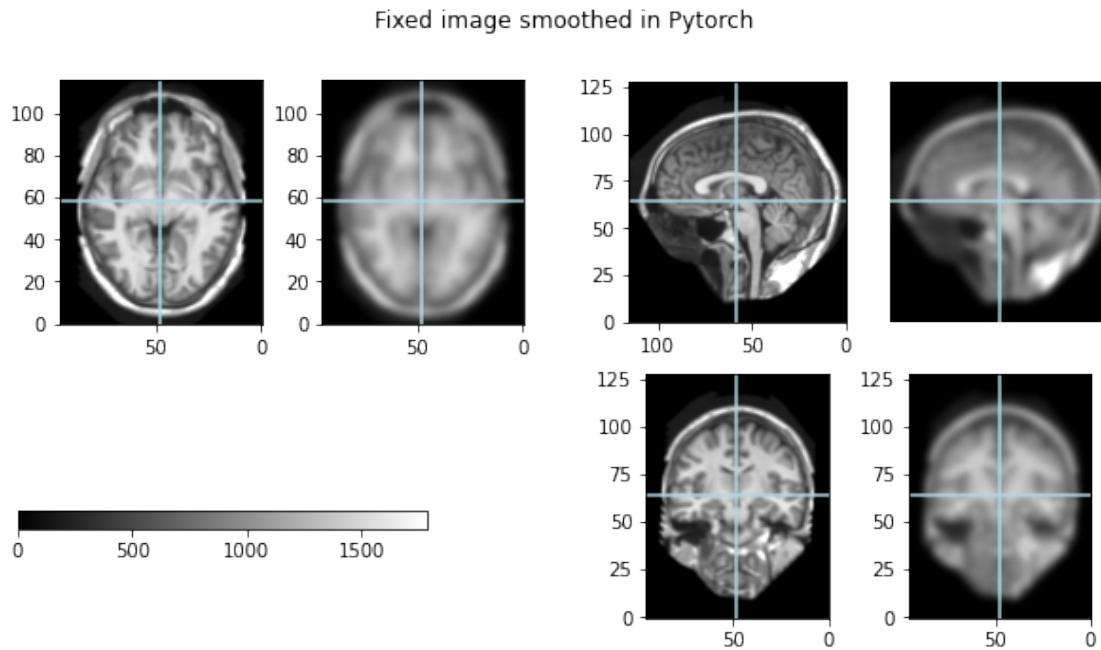
```

Here is an example of the fixed image smoothed in Pytorch

```

[14]: I_sm = my_pytorch_gaussian_lpf(I_fix, sigma=2.0)
my_view_multi([I_fix, I_sm], hdr_fix, title="Fixed image smoothed in Pytorch")

```



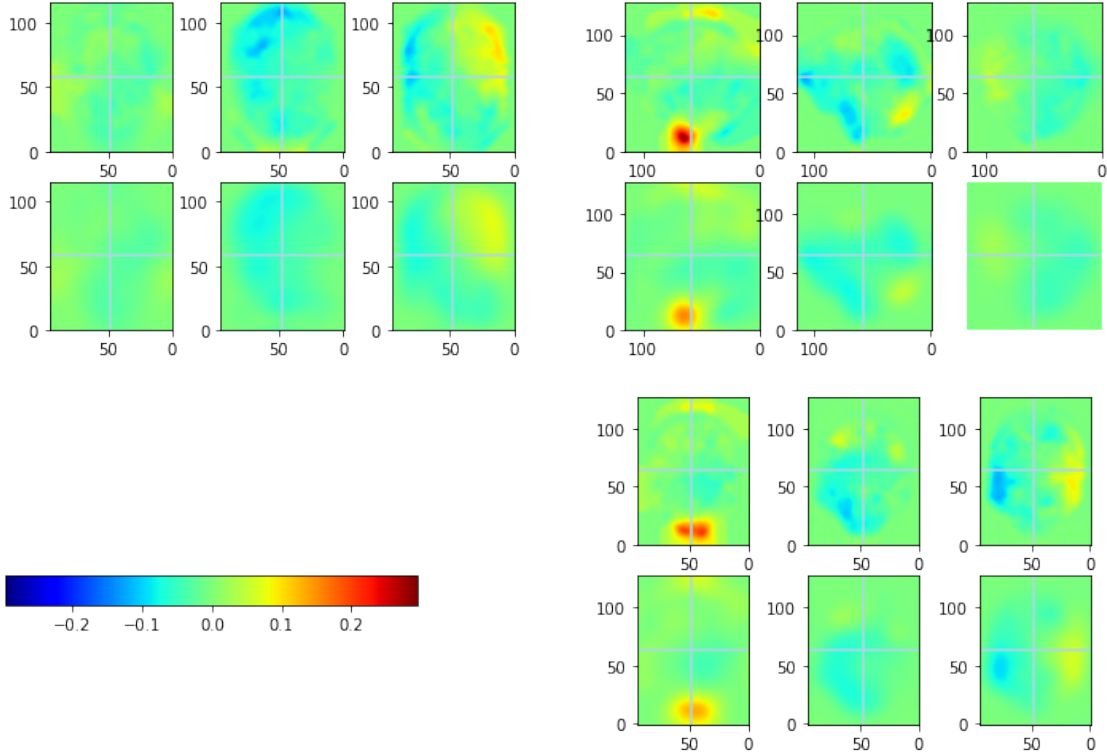
Here is an example of a deformation field smoothed in PyTorch

```

[15]: I_svf_sm = my_pytorch_gaussian_lpf_transform(I_svf, sigma=6.0)
my_view_multi(torch.cat((I_svf, I_svf_sm)), hdr_fix, is_warp=True,
               transpose=True, cmap='jet', figwidth=12,
               title='SVF (top) vs. Smoothed SVF (bottom)')

```


SVF (top) vs. Smoothed SVF (bottom)



3.5 Optical flow computation

Optical flow will be discussed in the Deformable Registration lectures and is covered in Section 3.2.1 of the NVR Chapter. Recall that optical flow infers a velocity field from a pair of images. When these images are frames from a video sequence, the optical flow is an estimate of how fast objects or particles are moving in the video. When doing inter-subject registration, optical flow does not have the same clear physical meaning, but it can be thought of as an infinitesimal displacement that, when applied to the moving image, makes it a little more similar to the fixed image.

3.5.1 Optical Flow Algorithms

Your task is to implement **two versions** of the optical flow approximation.

One of them will be a simple approximation that is obtained if you differentiate the mean square difference metric (MSD, also known as the sum of squared differences metric, or SSD) between $I(x)$ and $J^i(x + v(x))$ with respect to v . This approximation is referred to as `grad_msd` in the code below.

$$v = G_{\sigma} * (I - J^i) \nabla J^i$$

The other one can be one of the following: * The original *Demons* approximation, given in Algorithm

3 (line 2b) of the NVR Chapter (Section 3.2.2). * The Lukas-Kanade algorithm (advanced) * The Horn and Schunk algorithm (advanced)

3.5.2 Implementation as a Python Class

When working with PyTorch to train neural networks we typically not only write Python functions, but also [Python classes](#). A class is an object-oriented programming concept that allows data and functions to be abstracted away into a single entity. We will write a Python class to capture the logic of optical flow computation. The class will have a constructor (`__init__` function) that users call to configure the class (e.g., specify which optical flow method to use or how much smoothing to perform) and a method called `forward` that will perform the actual computation. Once we create an instance of this class, we can call `forward` multiple times with different image inputs.

TASK 5 (20 pts): complete the class definition below by writing the code for the `forward` function

Notes/Hints: * You may find the function `torch.where` useful when implementing the conditional expression in the Demons formula * Other functions that you will need here are `torch.square`, `torch.sum` and `torch.permute` * For the `grad_msd` method, the magnitude of the velocity field returned by `forward` will be very large (on the order of 10^7). This is okay because during registration we will normalize the velocity field so that the maximum displacement is on the order of one voxel.

```
[16]: class MyOpticalFlow:
    """A class that implements multiple optical flow algorithms"""

    def __init__(self, I, method='demons', sigma=2.0, demons_eps=0.001):
        """
        Initialize the optical flow computation

        :param I: Fixed image, represented as [N,1,D,H,W] tensor
        :param method:
            String, specifies which method is used for the computation.
        ↪ Available
            options are 'demons' (default) and 'grad_msd'
        :param sigma:
            Standard deviation for Gaussian smoothing in 'demons' and
        ↪ 'grad_msd'
            methods (passed to `my_gaussian_fft_kernel`)
        :param demons_eps:
            Only used if `method=='demons'`, sets the threshold on the
        ↪ denominator
            below which the flow is set to zero. See Algorithm 2 step 2b
        """
        self.method = method
        self.demons_eps = demons_eps
        self.kernel = my_gaussian_fft_kernel(I, sigma)

    def forward(self, I, J):
        """
```

```

    Compute optical flow for a given pair of fixed/moving images

    :param I: fixed image, represented as [N,1,D,H,W] tensor
    :param J: moving image, represented as [N,1,D,H,W] tensor
    :returns: optical flow velocity field, represented as a [N,D,H,W,3]
    ↪ tensor
    """
    if self.method=="grad_msdc":
        vbg = (I-J) * my_image_gradient(J) # v before Gaussian

        vbg = torch.permute(vbg, (0,2,3,4,1))

        v = my_pytorch_gaussian_lpf_transform(vbg, kernel=self.kernel)

        return v

    elif self.method=="demons":
        numerator = (I-J) * my_image_gradient(J)
        denominator = torch.sum(torch.square(my_image_gradient(J)), dim=1)\
            + torch.sum(torch.square(J-I), dim=1)

        vbg = numerator/denominator
        vbg = torch.where(denominator>self.demons_eps, vbg, torch.zeros(vbg.
    ↪shape, device=cuda))

        vbg = torch.permute(vbg, (0,2,3,4,1))

        v = my_pytorch_gaussian_lpf_transform(vbg, kernel=self.kernel).
    ↪float()

        return v

```

3.5.3 Check Optical Flow Computation

Visualize the optical flow fields that you obtain for the two methods using `my_view_multi` using code below. The output should look like `files/oflow_example.png`.

```

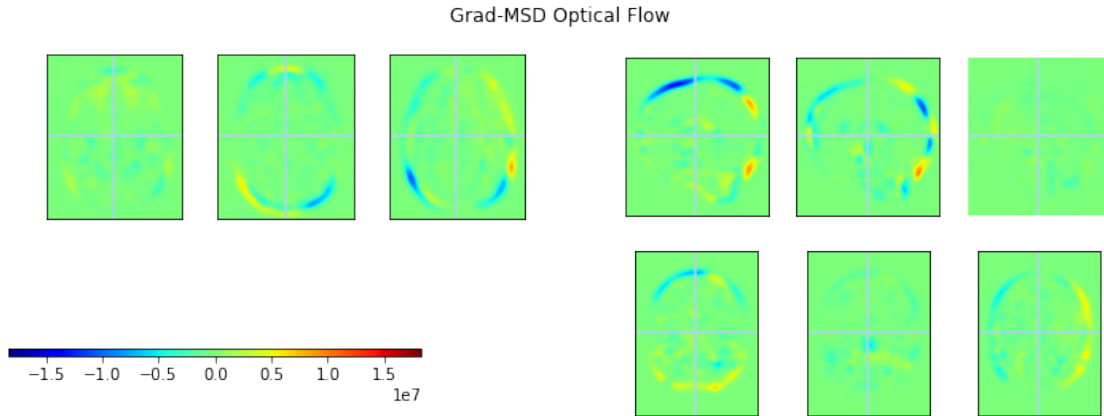
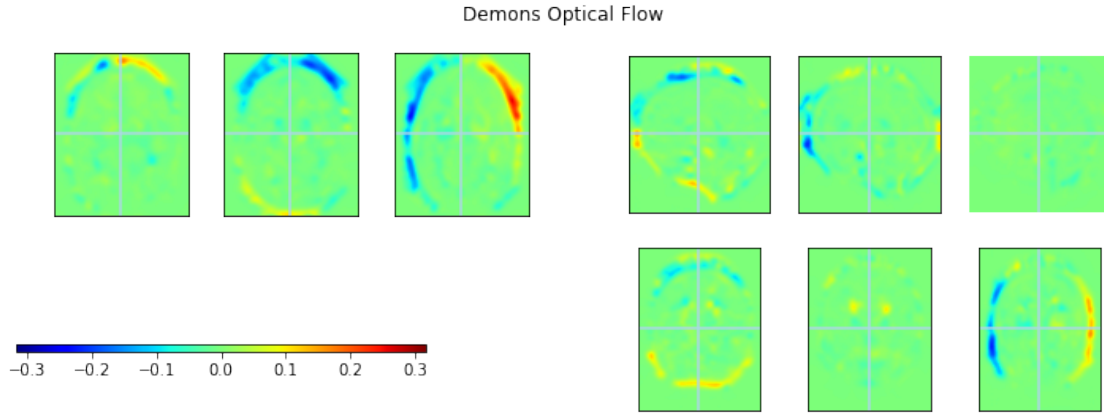
[17]: of_demons = MyOpticalFlow(I_fix, method='demons', sigma=2.0)
      v_demons = of_demons.forward(I_fix, I_mov)

      of_gradmsdc = MyOpticalFlow(I_fix, method='grad_msdc', sigma=2.0)
      v_gradmsdc = of_gradmsdc.forward(I_fix, I_mov)

      my_view_multi(v_demons, hdr_fix, cmap='jet', is_warp=True, transpose=True,
                    figwidth=12, axis_visible=False, title='Demons Optical Flow')

```

```
my_view_multi(v_gradmsd, hdr_fix, cmap='jet', is_warp=True, transpose=True,
              figwidth=12, axis_visible=False, title='Grad-MSD Optical Flow')
```



3.6 Log-Demons Algorithm

3.6.1 Full-Resolution Method

Finally, we are ready to implement the log-Demons algorithm, described in Algorithm 4 in the NVR chapter (Section 3.2.3). Again, we will implement it as a Python class `MyLogDemonsAlgorithm`

The method `__init__` will be responsible for storing the variables needed by the class (i.e., lines of code line `self.I = I`) and initializing the sampling grid and other objects that are repeatedly used during the algorithm.

The method `step` is the portion of the algorithm that is iterated. At every call to `step`, the spatial transformation ϕ^i and the stationary velocity field u^i are updated. Note that consistently with how we implemented scaling and squaring, we use the displacement field w^i to represent the spatial transformation ϕ^i in our code. The two are related by $\phi^i(x) = x + w^i(x)$.

The code for MyLogDemonsAlgorithm is provided to you, but it requires that gradient computation and scaling and squaring algorithms be implemented correctly.

```
[18]: class MyLogDemonsAlgorithm:
    """A class that performs log-Demons registration between two images"""

    def __init__(self, I, J, of, tau=0.5, eps_prime = 0.5):
        """
        Initialize the Log Demons algorithm

        :param I: Fixed image, represented as an [N,1,D,H,W] tensor
        :param J: Moving image, represented as an [N,1,D,H,W] tensor
        :param of:
            Instance of class MyOpticalFlow used to perform optical flow
        ↪computations
        :param tau:
            Standard deviation for 'diffusion-like' regularization (Algorithm
        ↪4, Line 2c)
        :param eps_prime:
            Normalization factor. The velocity field computed by optical flow
        ↪will
            be scaled so that the longest velocity vector is equal to eps_prime
            voxel widths. See second bullet point after Algorithm 3.
        """
        self.I = I
        self.J = J
        self.of = of
        self.kernel_tau = my_gaussian_fft_kernel(I, tau)
        self.eps_prime = eps_prime
        self.grid = my_create_pytorch_grid(I.shape, dtype=I.dtype, device=I.
        ↪device)

    def step(self, w_i, u_i):
        """
        Execute one iteration of the Log Demons algorithm (Algorithm 4, Line 2)

        :param w_i:
            Displacement field corresponding to the current estimate of the
        ↪spatial
            transformation. Transformation  $\phi_i$  is given by  $\phi_i(x) = x + w_i(x)$ .
        ↪+  $w_i(x)$ .
            Represented as an [N,1,D,H,W] tensor.
        :param u_i:
            Current estimate of the stationary velocity field  $u_i$ 
            Represented as an [N,1,D,H,W] tensor.

        :returns:
```

```

        Tuple (Ji, rms, w_i1, u_i1) where

        - `Ji` is the moving image transformed by the spatial_
        ↪ transformation  $\phi_i$ 
        - `msd` is the mean square intensity difference between I and  $J^i$ 
        - `w_i1` is the updated displacement field corresponding to_
        ↪  $\phi_{i+1}$ 
        - `u_i1` is the updated stationary velocity field
        """

        # Resample the moving image using current displacement field w_i
        J_i = tfun.grid_sample(self.J, self.grid + w_i, align_corners=False)

        # Compute the RMS metric between fixed image I and resampled moving_
        ↪ image
        msd = torch.mean((self.I - J_i).square())

        # Compute the optical flow field between fixed image I and resampled_
        ↪ moving image
        v_i = self.of.forward(self.I, J_i)

        # Scale the optical flow field so the longest vector does not exceed_
        ↪ eps_prime
        # (eps_prime is specified in voxel units, e.g., 0.5 voxel).
        scale = self.eps_prime * np.min(2.0 / np.array(self.I.shape[2:5])) /_
        ↪ torch.max(v_i)

        # Compute the updated SVF using the approximation of the Lie bracket_
        ↪ term in Algorithm 4 line 2c
        u_i1 = my_pytorch_gaussian_lpf_transform(u_i + scale * v_i, kernel=self._
        ↪ kernel_tau)

        # Perform scaling and squaring to get a new transformation
        w_i1 = my_scaling_and_squaring(u_i1, self.grid)

        # Return the updated w and u
        return J_i, msd, w_i1, u_i1

```

The code below should be used to test out your algorithm. It should produce output similar to what is shown in files/reg_example.png. The objective function should start at ~96000 and go down to ~13000. It might take a minute or two on a CPU to complete the 200 iterations.

```

[19]: log_demons = MyLogDemonsAlgorithm(I_fix, I_mov, of_gradmsd)
w_i = torch.zeros_like(I_fix).repeat((1,3,1,1,1)).permute(0,2,3,4,1)
u_i = w_i.clone()
for k in range(200):
    J_i, msd, w_i, u_i = log_demons.step(w_i, u_i)

```

```

if k % 20 == 0:
    print('Iter %03d, MSD = %8.4f' % (k, msd))

```

```

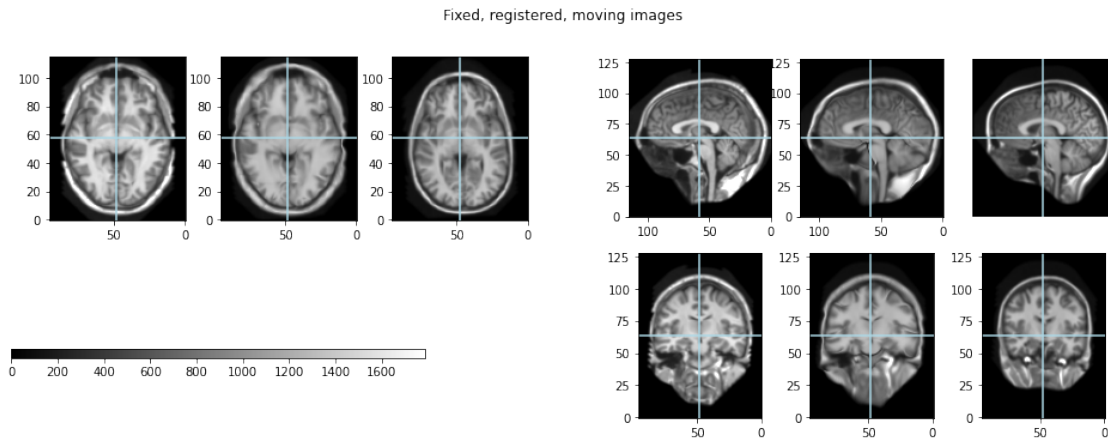
Iter 000, MSD = 96846.3906
Iter 020, MSD = 52207.1992
Iter 040, MSD = 31559.2520
Iter 060, MSD = 22472.2324
Iter 080, MSD = 18530.4863
Iter 100, MSD = 15748.2754
Iter 120, MSD = 14349.0088
Iter 140, MSD = 14096.1768
Iter 160, MSD = 13312.1777
Iter 180, MSD = 13139.9023

```

```

[20]: my_view_multi([I_fix, J_i, I_mov], hdr_fix, figwidth=16,
                    title='Fixed, registered, moving images')

```



3.6.2 Multi-Resolution Log Demons

A major difference between the registration method above and tools like ANTS is that the latter use multi-resolution schemes for significantly faster optimization. For example, when registering images of 256x256x256 voxels, one would typically subsample to 64x64x64 and run registration at that resolution; then use the result to initialize registration at 128x128x128 resolution, and run that for some number of iterations and use the result to initialize full-resolution registration. This has a dramatic effect on registration quality. This section provides code for multi-resolution Log Demons registration. There is no code for you to write in this section.

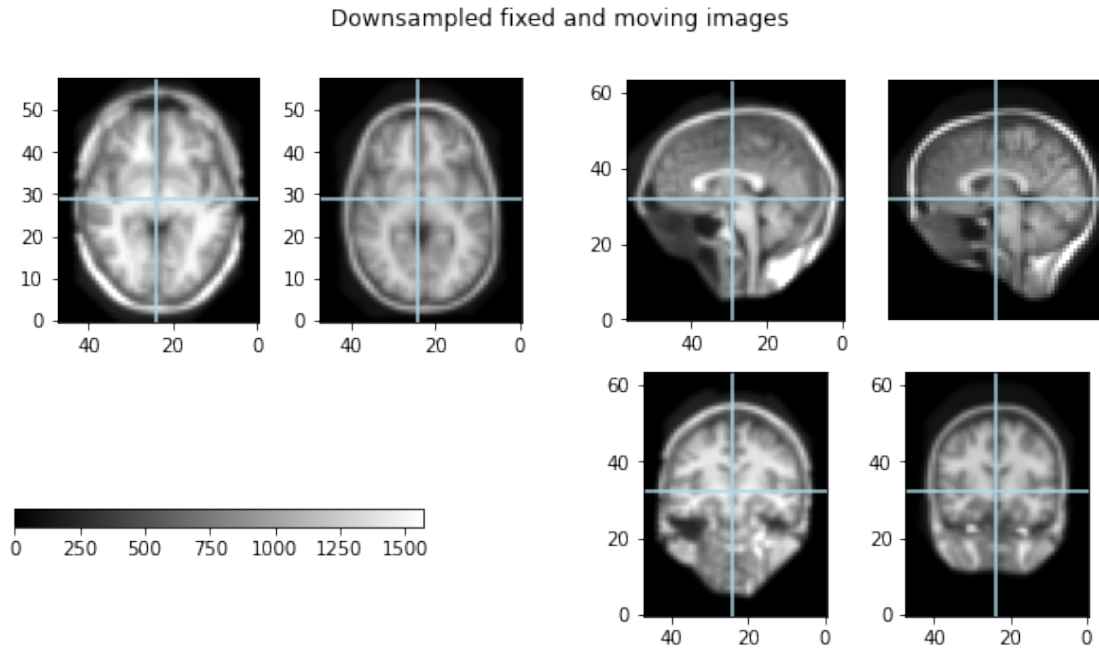
First, we coded a function to downsample an image represented as a PyTorch tensor (`my_image_downsample` in `be537_hw2`).

```

[21]: I_fix_ds = my_image_downsample(I_fix, 2.0)
      I_mov_ds = my_image_downsample(I_mov, 2.0)

```

```
hdr_fix_ds = my_adjust_nifti_header_for_resample(hdr_fix, I_fix_ds.shape[2:5])
my_view_multi([I_fix_ds, I_mov_ds], hdr_fix_ds,
              title="Downsampled fixed and moving images")
```



This function runs the log-Demons algorithm in a multi-resolution setting

```
[22]: def my_log_demons_multires(I_fix, I_mov, iter_vec,
                                sigma=2.0, tau=0.5, eps_prime=0.5,
                                of_method='grad_msd', of_param = {},
                                silent=False):

    """
    Multi-resolution log Demons algorithm

    :param I_fix: Fixed image, represented as an [N,1,D,H,W] tensor
    :param I_mov: Moving image, represented as an [N,1,D,H,W] tensor
    :param iter_vec:
        Number of iterations at each resolution level, e.g., [100,40] will
        do 100 iterations at 2x downsampling, 40 at full resolution
    :param sigma: See MyOpticalFlow
    :param tau: See MyLogDemonsAlgorithm
    :param eps_prime: See MyLogDemonsAlgorithm
    :param of_method: See MyOpticalFlow
    :param of_param: See MyOpticalFlow

    :returns:
        Tuple (I_res, rms, w, u) where
```



```

- `I_res` is the moving image transformed by the spatial transformation_
↪  $\phi$ 
- `rms` is the root mean square intensity difference between  $I$  and  $I_{res}$ 
- `w` is the displacement field corresponding to  $\phi$ 
- `u` is the stationary velocity field corresponding to  $\phi$ 
"""

# Current u
u = None

# Iterate over resolution levels
for level, n_iter in enumerate(iter_vec):

    # Downsample images by current factor
    factor = 2 ** (len(iter_vec) - level - 1)
    I_fix_ds = my_image_downsample(I_fix, factor)
    I_mov_ds = my_image_downsample(I_mov, factor)
    dim_ds = I_fix_ds.shape[2:5]

    # Resample u to the current iteration
    if u is not None:
        u = tfun.interpolate(u.permute(0,4,1,2,3),
                             tuple(dim_ds)).permute(0,2,3,4,1)
    else:
        u = torch.zeros_like(I_fix_ds).permute(0,2,3,4,1).repeat(1,1,1,1,3)

    # Create optical flow for this iteration
    of = MyOpticalFlow(I_fix_ds, sigma=sigma, method=of_method, **of_param)

    # Create the algorithm for this iteration
    log_demons = MyLogDemonsAlgorithm(I_fix_ds, I_mov_ds, of, tau=tau,
                                       eps_prime=eps_prime)

    # Compute the displacement field w
    w = my_scaling_and_squaring(u, log_demons.grid)

    # Perform iterations
    for k in range(n_iter):
        I_res, msd, w, u = log_demons.step(w, u)
        if k % 10 == 0:
            if not silent:
                print('Level %01d  Iter %03d  RMS %8.4f' % (level, k,
                                                            msd.item()))

    return I_res, msd, w, u

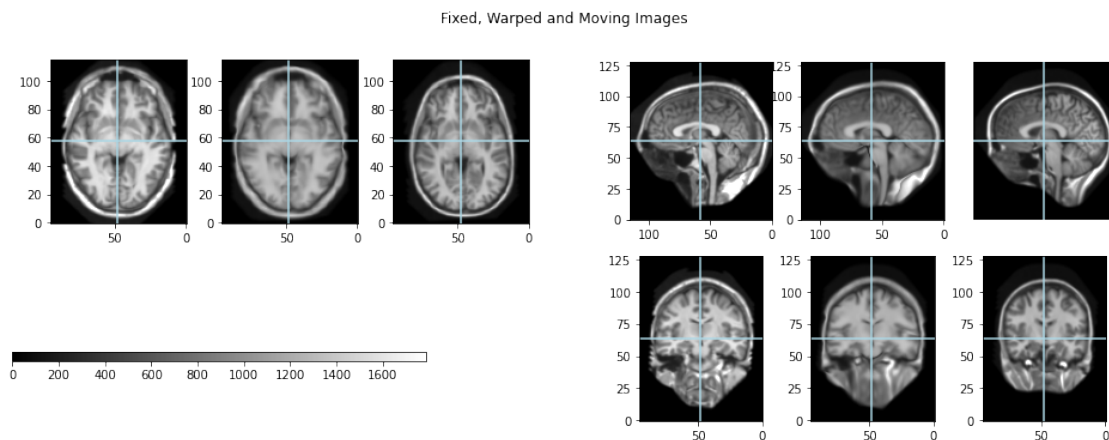
```

Run a multi-resolution registration. This registration should run faster than the one above and get the final RMS that is under 13,000.

```
[23]: I_res, rms, phi, u = my_log_demons_multires(I_fix, I_mov, [80,80,40],
                                             of_method='grad_msd', sigma=1.0)
```

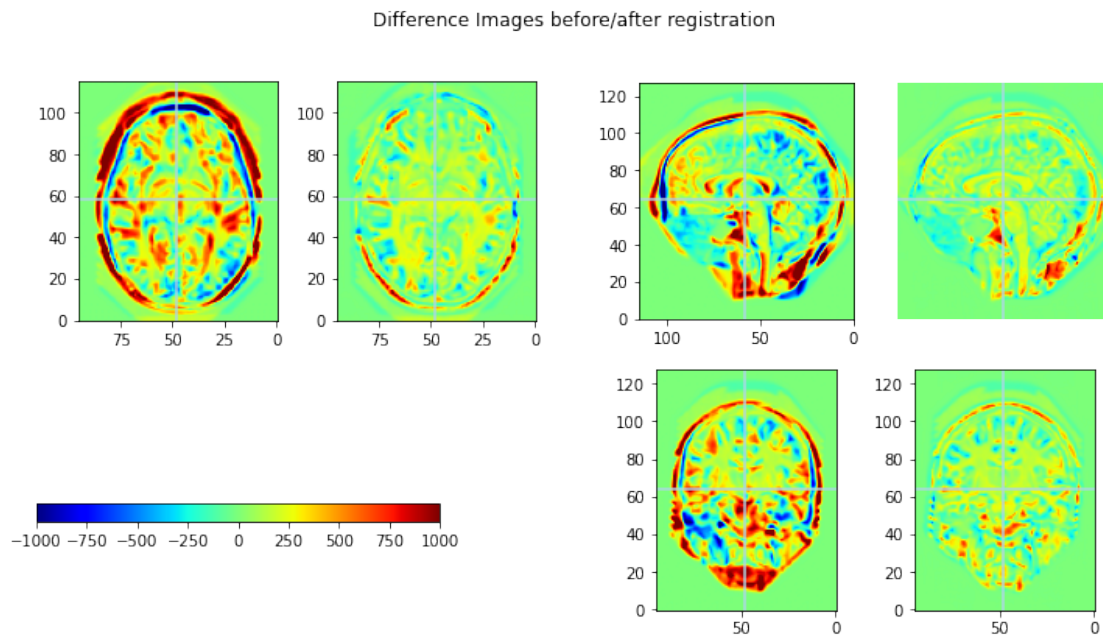
```
Level 0 Iter 000 RMS 38671.7969
Level 0 Iter 010 RMS 25605.3965
Level 0 Iter 020 RMS 23156.2090
Level 0 Iter 030 RMS 20229.6172
Level 0 Iter 040 RMS 19559.4141
Level 0 Iter 050 RMS 19037.0918
Level 0 Iter 060 RMS 18722.2539
Level 0 Iter 070 RMS 18559.9395
Level 1 Iter 000 RMS 36157.4648
Level 1 Iter 010 RMS 18300.3711
Level 1 Iter 020 RMS 13684.6992
Level 1 Iter 030 RMS 11146.5049
Level 1 Iter 040 RMS 10658.3428
Level 1 Iter 050 RMS 10496.4746
Level 1 Iter 060 RMS 10224.4717
Level 1 Iter 070 RMS 10139.5508
Level 2 Iter 000 RMS 16456.1543
Level 2 Iter 010 RMS 13733.7529
Level 2 Iter 020 RMS 13152.0605
Level 2 Iter 030 RMS 12773.3506
```

```
[24]: my_view_multi([I_fix,I_res,I_mov], hdr_fix, figwidth=16,
                    title='Fixed, Warped and Moving Images')
```



Plot difference images before and after registration

```
[25]: my_view_multi([I_fix-I_mov,I_fix-I_res], hdr_fix, cmap='jet',
                    crange=[-1000,1000], figwidth=12,
                    title='Difference Images before/after registration')
```



3.7 Compare Affine and Deformable Registration for 10 Subject Pairs

In this final sub-task of the base component of the assignment, you will perform deformable registration between 10 randomly selected pairs of subjects in the dataset and evaluate how well these registrations align anatomical labels between images compared to affine registration. To make things simpler, we computed the affine transformations between all image pairs ahead of time.

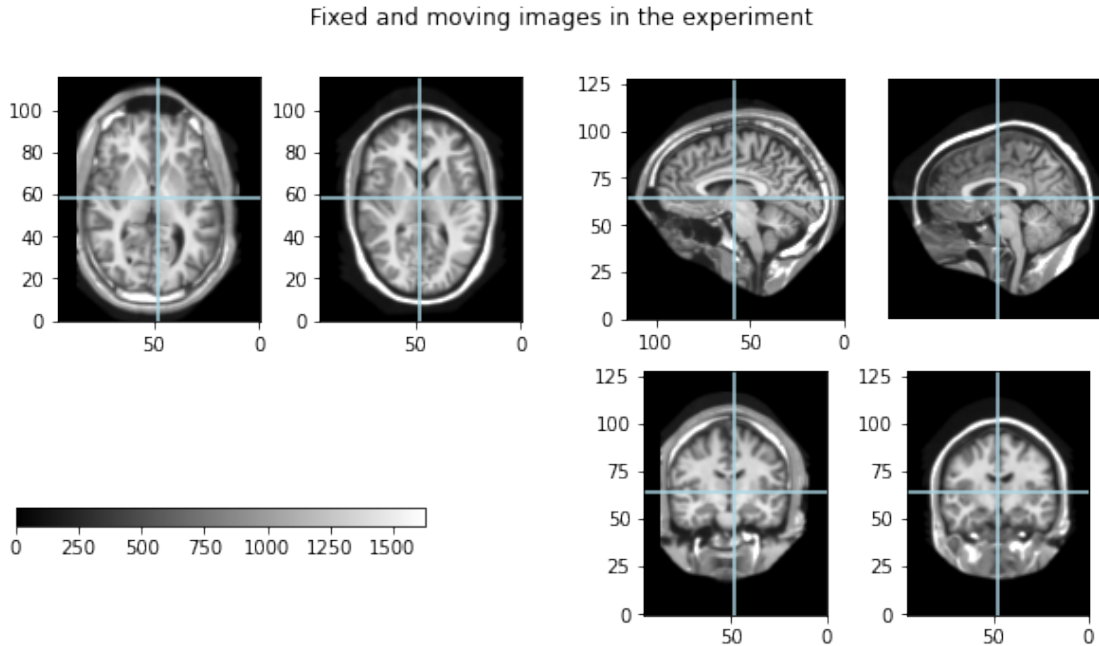
3.7.1 Class to Load Registration Experiments

We also provide some helper code to save you some tedious work. The class `RegistrationExperiment` provided in the `be537hw2.py` lets you easily load and access all the data associated with each registration experiment. This class provides access to the following variables:

- `exp.I_fix`: Fixed image, represented as a PyTorch tensor
- `exp.S_fix`: Fixed image segmentation
- `exp.I_mov`: Moving image
- `exp.S_mov`: Moving image segmentation
- `exp.hdr_fix`: Fixed image header (for calling `my_view`)
- `exp.grid`: Fixed image coordinate grid (for calling `grid_sample`)
- `exp.A`, `exp.b`: Affine transform used to map the moving image into fixed image space.

The code below loads a registration experiment with subject 1002 as fixed and 1014 as moving.

```
[26]: exp = RegistrationExperiment('1002','1014', data_path, dtype=torch.float32,
                                   device=cuda)
my_view_multi([exp.I_fix, exp.I_mov], exp.hdr_fix,
               title='Fixed and moving images in the experiment')
```




3.7.2 Function to apply affine and deformable transformations jointly

Another useful function we provided is `my_transform_affine_and_warp`. It allows you to apply the a combination of affine and deformable transformations to a moving image. If the deformable transformation is not specified (as in the line below), only the affine transformation is applied.

```
[27]: I_aff = my_transform_affine_and_warp(exp.I_mov, exp.grid, exp.A, exp.b)
my_view_multi([exp.I_fix, I_aff], exp.hdr_fix,
               title='Fixed and moving images in the experiment after affine_
               ↪transform')
```

Fixed and moving images in the experiment after affine transform




Now that we have the fixed image `exp.I_fix` and the affine-transformed moving image `I_aff`, we can apply your log-Demons algorithm to perform deformable registration.

```
[28]: _, _, w_opt, _ = my_log_demons_multires(exp.I_fix, I_aff, [80,80,40],
                                             of_method='grad_msd', sigma=1.0,
                                             silent=True)
```

To apply the combination of deformable and affine transformations to the moving image, we use the command below. This helps avoid double interpolation and associated aliasing, that would happen if we first applied the affine transform and then separately applied the deformable transformation

```
[29]: I_def = my_transform_affine_and_warp(exp.I_mov, exp.grid, exp.A, exp.b, w_opt)
my_view_multi([exp.I_fix, I_def], exp.hdr_fix,
              title='Fixed and moving images after deformable transformation')
```

Fixed and moving images after deformable transformation




3.7.3 Applying transformations to segmentation images

In order to apply your affine and deformable transformations to the anatomical segmentation of the moving image, we can use the following code, which uses nearest neighbor interpolation to avoid interpolating linearly between adjacent anatomical labels.

```
[30]: S_aff = my_transform_affine_and_warp(exp.S_mov, exp.grid, exp.A, exp.b,
                                         mode='nearest')
S_def = my_transform_affine_and_warp(exp.S_mov, exp.grid, exp.A, exp.b, w_opt,
                                     mode='nearest')
my_view_multi([exp.S_fix, exp.S_mov, S_aff, S_def], exp.hdr_fix,
              cmap='jet', layout='A;S',
              title='Segmentations: fixed, moving, moving after affine, moving_
↳after deformable.')
```

Segmentations: fixed, moving, moving after affine, moving after deformable.



3.7.4 Measuring alignment between segmentations

Finally, to measure how well the registration aligned anatomical structures between the fixed and moving images, use the function `my_generalized_dice` to measure the [generalized Dice similarity coefficient](#) between the segmentations. GDSC is a measure between 0 and 1 with 0 indicating no overlap and 1 indicating perfect overlap. The code below computes GDSC at different stages of the registration.

```
[31]: print('Alignment between segmentations images %s and %s' % (exp.fixed_id, exp.
      ↪moving_id))
print('  GDSC before registration: ', my_generalized_dice(exp.S_fix, exp.S_mov))
print('  GDSC after affine registration', my_generalized_dice(exp.S_fix, S_aff))
print('  GDSC after deformable registration', my_generalized_dice(exp.S_fix,
      ↪S_def))
```

Alignment between segmentations images 1002 and 1014

GDSC before registration: tensor(0.1664, device='cuda:0')

GDSC after affine registration tensor(0.3797, device='cuda:0')

GDSC after deformable registration tensor(0.4727, device='cuda:0')

3.7.5 Evaluate GDSC for 10 subject pairs

```
[32]: def evaluate_gdsc(id_fix, id_mov, grad_msd=False, demons=False):
      """_summary_

      :param id_fix: ID of the fixed image
      :param id_mov: ID of the moving image
```

```

        :param grad_msd: whether to do grad_msd deformable registration, defaults_
        ↪to False
        :param demons: whether to do demons deformable registration, defaults to_
        ↪False

        :returns:
            Tuple (gdsc_all, segmentations) where

            - `gdsc_all` is the list of GDSC for none, affine, deformable_
            ↪registration
            - `segmentations` is the list of affine and deformable registered_
            ↪segmentations
        """

        gdsc_all = []
        segmentations = []

        exp = RegistrationExperiment(str(id_fix), str(id_mov), data_path,
                                    dtype=torch.float32, device=cuda)

        gdsc_no_reg = my_generalized_dice(exp.S_fix, exp.S_mov).cpu().numpy()
        gdsc_all.append(gdsc_no_reg)

        print(' GDSC before registration: ', gdsc_no_reg)

        # get affine registered image
        I_aff = my_transform_affine_and_warp(exp.I_mov, exp.grid, exp.A, exp.b)
        S_aff = my_transform_affine_and_warp(exp.S_mov, exp.grid, exp.A, exp.b,
                                             mode='nearest')
        gdsc_affine = my_generalized_dice(exp.S_fix, S_aff).cpu().numpy()

        segmentations.append(S_aff)
        gdsc_all.append(gdsc_affine)

        print(' GDSC after affine registration', gdsc_affine)

        if grad_msd:
            # grad msd deformable registration
            _, _, w_gradmsd, _ = my_log_demons_multires(exp.I_fix, I_aff,
                                                         [80,80,40],
                                                         of_method='grad_msd',
                                                         sigma=1.0, silent=True)
            S_gradmsd = my_transform_affine_and_warp(exp.S_mov, exp.grid, exp.A,
                                                         exp.b, w_gradmsd,
                                                         mode='nearest')
            gdsc_def_gradmsd = my_generalized_dice(exp.S_fix, S_gradmsd).cpu().
            ↪numpy()

```



```

segmentations.append(S_gradmsd)
gdsc_all.append(gdsc_def_gradmsd)

print(' GDSC after grad_msd deformable registration', gdsc_def_gradmsd)

if demons:
    # demons deformable registration
    _, _, w_demons, _ = my_log_demons_multires(exp.I_fix, I_aff, [80,80,40],
                                                of_method='demons',
                                                sigma=1.0, silent=True)
    S_demons = my_transform_affine_and_warp(exp.S_mov, exp.grid, exp.A,
                                            exp.b, w_demons, mode='nearest')
    gdsc_def_demons = my_generalized_dice(exp.S_fix, S_demons).cpu().numpy()
    segmentations.append(S_demons)
    gdsc_all.append(gdsc_def_demons)

print(' GDSC after demons deformable registration', gdsc_def_demons)

return gdsc_all, segmentations

```

```

[33]: def print_mean_std(my_list):
    my_array = np.array(my_list)
    mean = np.mean(my_array)
    std = np.std(my_array)

    return mean, std

```

TASK 6 (10 points): Pick 10 random pairs of fixed/moving subjects, perform registration, and record GDSC before registration, after affine registration, after deformable registration using Demons method, and after deformable registration using one of the other optical flow methods you implemented. Report summary statistics and conclude which technique performed best.

Method	Average GDSC
No registration	0.2461 \pm 0.0743
Affine only	0.3769 \pm 0.0328
Affine + Grad MSD	0.4940 \pm 0.0224
Affine + Demons	0.4463 \pm 0.0312

```

[34]: list_atlas = [1000, 1001, 1002, 1006, 1007, 1008, 1009, 1010, 1011, 1012, 1013,
                  1014, 1015, 1017, 1036]

gdsc_no_reg = []
gdsc_affine = []
gdsc_gradmsd = []
gdsc_demons = []

```

```

for i in range(10):
    # pick two IDs randomly from list above
    id_fix, id_mov = np.random.choice(list_atlas, 2, replace=False)
    print(id_fix, id_mov)

    gdsc_all, _ = evaluate_gdsc(id_fix, id_mov, grad_msd=True, demons=True)

    gdsc_no_reg.append(gdsc_all[0])
    gdsc_affine.append(gdsc_all[1])
    gdsc_gradmsd.append(gdsc_all[2])
    gdsc_demons.append(gdsc_all[3])

```

1013 1036

GDSC before registration: 0.22458068
GDSC after affine registration 0.41836646
GDSC after grad_msd deformable registration 0.5040101
GDSC after demons deformable registration 0.44491392

1014 1036

GDSC before registration: 0.13437201
GDSC after affine registration 0.32930702
GDSC after grad_msd deformable registration 0.47471428
GDSC after demons deformable registration 0.4444901

1000 1006

GDSC before registration: 0.16061912
GDSC after affine registration 0.34549358
GDSC after grad_msd deformable registration 0.4864284
GDSC after demons deformable registration 0.4590709

1009 1010

GDSC before registration: 0.34810483
GDSC after affine registration 0.3859839
GDSC after grad_msd deformable registration 0.5215652
GDSC after demons deformable registration 0.4297017

1001 1014

GDSC before registration: 0.29346314
GDSC after affine registration 0.367543
GDSC after grad_msd deformable registration 0.48808372
GDSC after demons deformable registration 0.45490834

1006 1000

GDSC before registration: 0.16061912
GDSC after affine registration 0.34566307
GDSC after grad_msd deformable registration 0.4872934
GDSC after demons deformable registration 0.42142546

1013 1007

GDSC before registration: 0.22376765
GDSC after affine registration 0.35946524
GDSC after grad_msd deformable registration 0.49983105

```

GDSC after demons deformable registration 0.4653676
1008 1011
GDSC before registration: 0.2757263
GDSC after affine registration 0.43848667
GDSC after grad_msd deformable registration 0.5365502
GDSC after demons deformable registration 0.51401544
1008 1009
GDSC before registration: 0.3564495
GDSC after affine registration 0.37907672
GDSC after grad_msd deformable registration 0.45155773
GDSC after demons deformable registration 0.38594896
1013 1006
GDSC before registration: 0.28333756
GDSC after affine registration 0.39988133
GDSC after grad_msd deformable registration 0.48959446
GDSC after demons deformable registration 0.44302535

```

```

[35]: print('GDSC before registration: ', print_mean_std(gdsc_no_reg))
      print('GDSC after affine registration', print_mean_std(gdsc_affine))
      print('GDSC after grad_msd deformable registration',
            print_mean_std(gdsc_gradmsd))
      print('GDSC after demons deformable registration', print_mean_std(gdsc_demons))

```

```

GDSC before registration: (0.246104, 0.07426206)
GDSC after affine registration (0.37692672, 0.032781854)
GDSC after grad_msd deformable registration (0.49396285, 0.022437815)
GDSC after demons deformable registration (0.44628677, 0.031206675)

```

4 Extension: Multi-Atlas Segmentation

```

[36]: list_atlas = [1000, 1001, 1002, 1006, 1007, 1008, 1009, 1010, 1011, 1012, 1013,
                  1014, 1015, 1017, 1036]

# randomly choose the 10 atlas images
atlas_imgs = np.random.choice(list_atlas, 10)

```

```

[37]: def get_strong_segmentation(id_fix, atlas_imgs=atlas_imgs):
      """Generate strong segmentation for a given image ID

      :param id_fix: ID for the target image
      :param atlas_imgs: list of IDs for the atlas images, defaults to atlas_imgs
      :return:
          Tuple (strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, exp.S_fix)
      where

      - `strong_seg_aff` is the affine registered strong segmentation
      - `strong_seg_def` is the deformable registered strong segmentation
      """

```

```

- `aff_gdsc` is the GDSC between the ground truth and affine registered,
↳strong segmentation
- `def_gdsc` is the GDSC between the ground truth and deformable,
↳registered strong segmentation
- `exp.S_fix` is the ground truth segmentation
"""
gdsc_no_reg = []
gdsc_affine = []
gdsc_def = []

weak_segmentations_aff = []
weak_segmentations_def = []

for id_mov in atlas_imgs:
    print(id_fix, id_mov)
    gdsc_all, segmentations = evaluate_gdsc(id_fix, id_mov, grad_msd=True)

    gdsc_no_reg.append(gdsc_all[0])
    gdsc_affine.append(gdsc_all[1])
    gdsc_def.append(gdsc_all[2])

    weak_segmentations_aff.append(segmentations[0])
    weak_segmentations_def.append(segmentations[1])

print('\nWEAK SEGMENTATIONS')
print(' Average GDSC before registration: ', print_mean_std(gdsc_no_reg))
print(' Average GDSC after affine registration: ',
↳print_mean_std(gdsc_affine))
print(' Average GDSC after deformable registration: ',
↳print_mean_std(gdsc_def))

strong_seg_aff = my_majority_vote(weak_segmentations_aff)
strong_seg_def = my_majority_vote(weak_segmentations_def)

exp = RegistrationExperiment(str(id_fix),str(id_mov), data_path,
                           dtype=torch.float32, device=cuda)


aff_gdsc = my_generalized_dice(exp.S_fix, strong_seg_aff).cpu().numpy()
def_gdsc = my_generalized_dice(exp.S_fix, strong_seg_def).cpu().numpy()

print('\nSTRONG SEGMENTATIONS')
print(' GDSC with only affine registration: ', aff_gdsc)
print(' GDSC with deformable registration: ', def_gdsc)

return strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, exp.S_fix

```

4.1 Target Image 1003

```
[38]: strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, S_fix =   
      ↪ get_strong_segmentation(1003)
```

```
1003 1000  
    GDSC before registration: 0.1316802  
    GDSC after affine registration 0.3823198  
    GDSC after grad_msd deformable registration 0.4354621  
1003 1036  
    GDSC before registration: 0.15730593  
    GDSC after affine registration 0.37248668  
    GDSC after grad_msd deformable registration 0.39088637  
1003 1001  
    GDSC before registration: 0.30236235  
    GDSC after affine registration 0.42282417  
    GDSC after grad_msd deformable registration 0.34501716  
1003 1001  
    GDSC before registration: 0.30236235  
    GDSC after affine registration 0.42282417  
    GDSC after grad_msd deformable registration 0.34501716  
1003 1036  
    GDSC before registration: 0.15730593  
    GDSC after affine registration 0.37248668  
    GDSC after grad_msd deformable registration 0.39088637  
1003 1014  
    GDSC before registration: 0.31679967  
    GDSC after affine registration 0.3829238  
    GDSC after grad_msd deformable registration 0.39848256  
1003 1014  
    GDSC before registration: 0.31679967  
    GDSC after affine registration 0.3829238  
    GDSC after grad_msd deformable registration 0.39848256  
1003 1009  
    GDSC before registration: 0.35558972  
    GDSC after affine registration 0.3873193  
    GDSC after grad_msd deformable registration 0.5135692  
1003 1007  
    GDSC before registration: 0.38080025  
    GDSC after affine registration 0.43082297  
    GDSC after grad_msd deformable registration 0.49873152  
1003 1012  
    GDSC before registration: 0.20964095  
    GDSC after affine registration 0.4415992  
    GDSC after grad_msd deformable registration 0.5050789
```

WEAK SEGMENTATIONS

```
Average GDSC before registration: (0.2630647, 0.08589055)
```

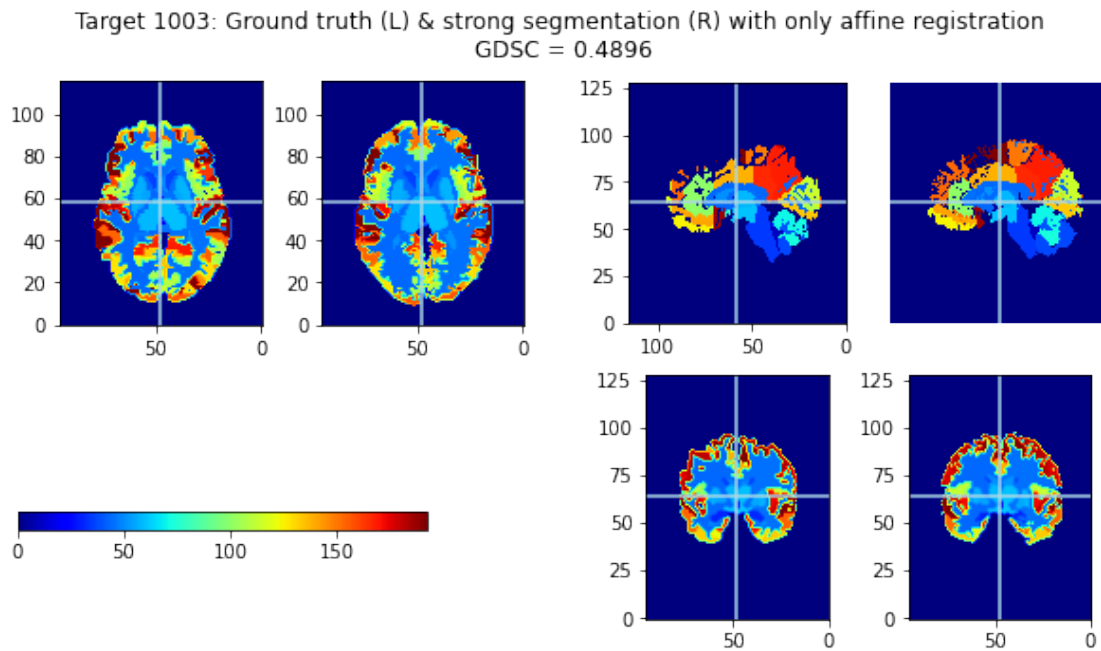
Average GDSC after affine registration: (0.39985305, 0.025085386)
Average GDSC after deformable registration: (0.4221614, 0.06020854)

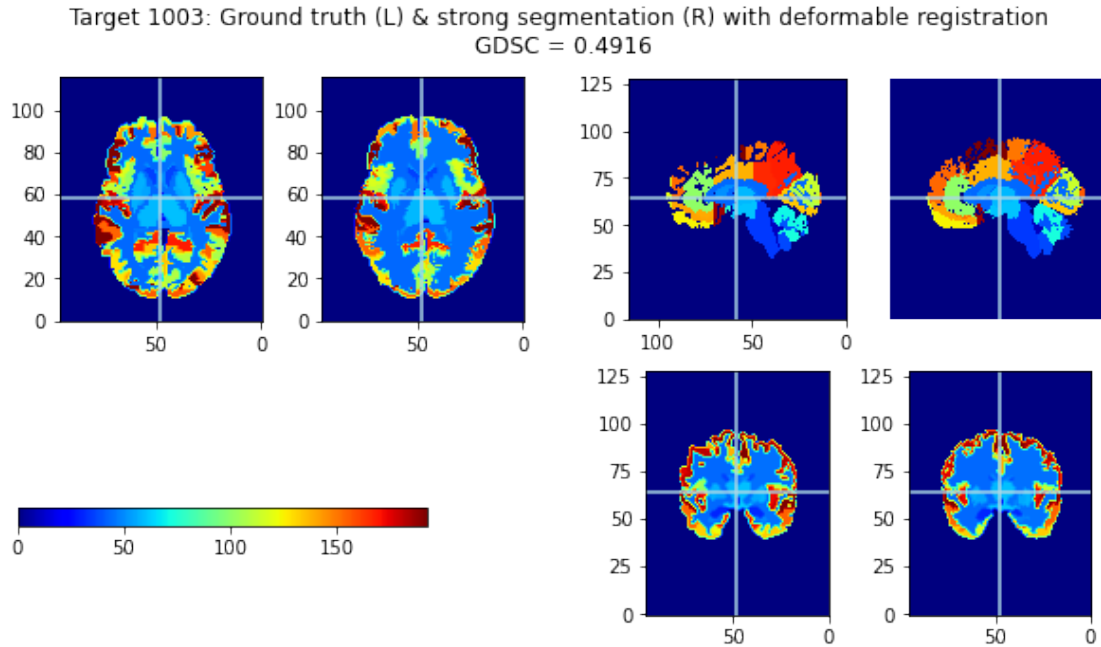
STRONG SEGMENTATIONS

GDSC with only affine registration: 0.4895567
GDSC with deformable registration: 0.4916204

```
[39]: my_view_multi([S_fix, strong_seg_aff], header=exp.hdr_fix, cmap='jet',
                    title=f"Target 1003: Ground truth (L) & strong segmentation (R)␣
                    ↳with only affine registration \nGDSC = {aff_gdsc:.4f}")

my_view_multi([S_fix, strong_seg_def], header=exp.hdr_fix, cmap='jet',
              title=f"Target 1003: Ground truth (L) & strong segmentation (R)␣
              ↳with deformable registration \nGDSC = {def_gdsc:.4f}")
```





4.2 Target Image 1004

```
[40]: strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, S_fix = get\_strong\_segmentation\(1004\)
```

```
1004 1000
GDSC before registration: 0.1809921
GDSC after affine registration 0.40627903
GDSC after grad_msd deformable registration 0.5095027
1004 1036
GDSC before registration: 0.22732563
GDSC after affine registration 0.3882557
GDSC after grad_msd deformable registration 0.5001966
1004 1001
GDSC before registration: 0.31397113
GDSC after affine registration 0.40873292
GDSC after grad_msd deformable registration 0.49865496
1004 1001
GDSC before registration: 0.31397113
GDSC after affine registration 0.40873292
GDSC after grad_msd deformable registration 0.49865496
1004 1036
GDSC before registration: 0.22732563
GDSC after affine registration 0.3882557
GDSC after grad_msd deformable registration 0.5001966
```

```

1004 1014
  GDSC before registration: 0.26888025
  GDSC after affine registration 0.37708142
  GDSC after grad_msd deformable registration 0.5004797
1004 1014
  GDSC before registration: 0.26888025
  GDSC after affine registration 0.37708142
  GDSC after grad_msd deformable registration 0.5004797
1004 1009
  GDSC before registration: 0.29996473
  GDSC after affine registration 0.36913952
  GDSC after grad_msd deformable registration 0.4326929
1004 1007
  GDSC before registration: 0.28079155
  GDSC after affine registration 0.38683254
  GDSC after grad_msd deformable registration 0.49758324
1004 1012
  GDSC before registration: 0.28710777
  GDSC after affine registration 0.44979498
  GDSC after grad_msd deformable registration 0.4808879

```

WEAK SEGMENTATIONS

```

Average GDSC before registration: (0.26692098, 0.040799823)
Average GDSC after affine registration: (0.3960186, 0.022282528)
Average GDSC after deformable registration: (0.49193293, 0.020844763)

```

STRONG SEGMENTATIONS

```

GDSC with only affine registration: 0.48550487
GDSC with deformable registration: 0.5952851

```

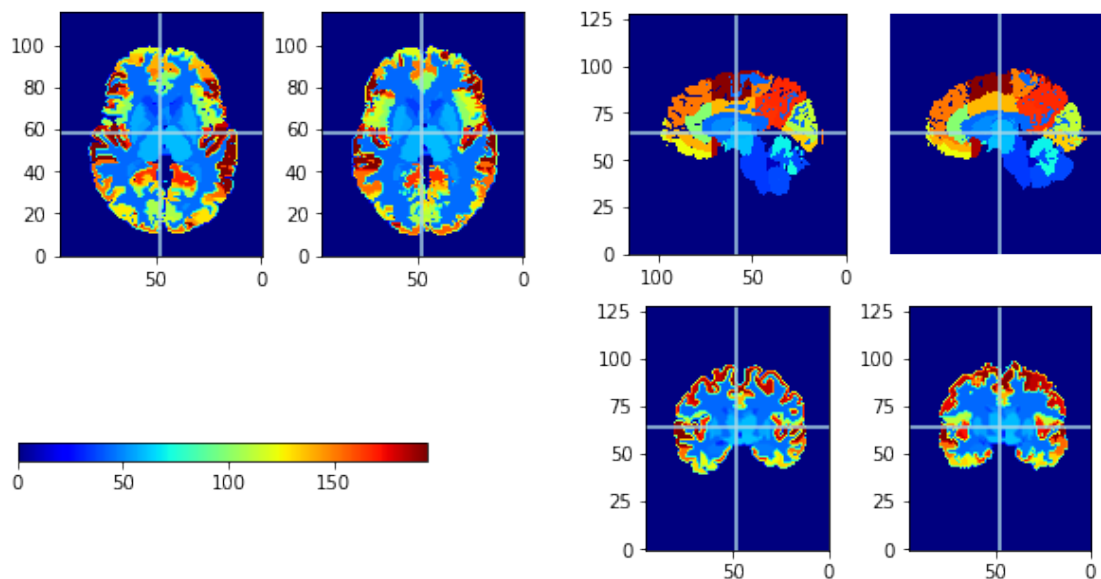
```

[41]: my_view_multi([S_fix, strong_seg_aff], header=exp.hdr_fix, cmap='jet',
                    title=f"Target 1004: Ground truth (L) & strong segmentation (R)␣
                    ↳with only affine registration \nGDSC = {aff_gdsc:.4f}")

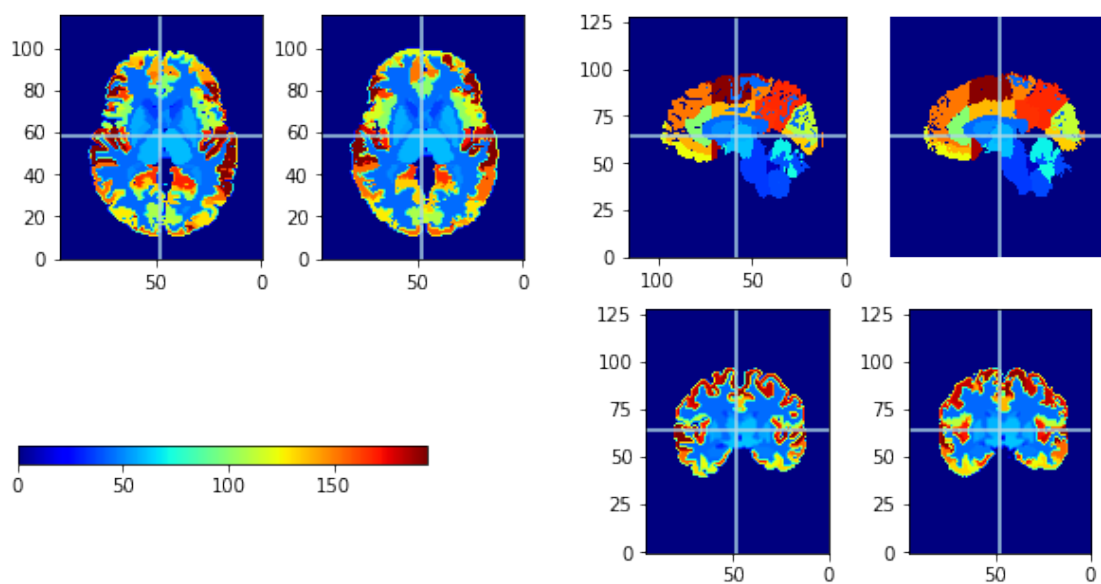
my_view_multi([S_fix, strong_seg_def], header=exp.hdr_fix, cmap='jet',
              title=f"Target 1004: Ground truth (L) & strong segmentation (R)␣
              ↳with deformable registration \nGDSC = {def_gdsc:.4f}")

```



Target 1004: Ground truth (L) & strong segmentation (R) with only affine registration
GDSC = 0.4855



Target 1004: Ground truth (L) & strong segmentation (R) with deformable registration
GDSC = 0.5953



4.3 Target Image 1005

```
[42]: strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, S_fix =   
      ↪ get_strong_segmentation(1005)
```

```
1005 1000  
    GDSC before registration: 0.20155624  
    GDSC after affine registration 0.37697098  
    GDSC after grad_msd deformable registration 0.4967663  
1005 1036  
    GDSC before registration: 0.2049951  
    GDSC after affine registration 0.36643764  
    GDSC after grad_msd deformable registration 0.4970466  
1005 1001  
    GDSC before registration: 0.28162566  
    GDSC after affine registration 0.36923528  
    GDSC after grad_msd deformable registration 0.46596432  
1005 1001  
    GDSC before registration: 0.28162566  
    GDSC after affine registration 0.36923528  
    GDSC after grad_msd deformable registration 0.46596432  
1005 1036  
    GDSC before registration: 0.2049951  
    GDSC after affine registration 0.36643764  
    GDSC after grad_msd deformable registration 0.4970466  
1005 1014  
    GDSC before registration: 0.28053266  
    GDSC after affine registration 0.39449316  
    GDSC after grad_msd deformable registration 0.48526144  
1005 1014  
    GDSC before registration: 0.28053266  
    GDSC after affine registration 0.39449316  
    GDSC after grad_msd deformable registration 0.48526144  
1005 1009  
    GDSC before registration: 0.2820854  
    GDSC after affine registration 0.34989637  
    GDSC after grad_msd deformable registration 0.42476666  
1005 1007  
    GDSC before registration: 0.28492218  
    GDSC after affine registration 0.35587177  
    GDSC after grad_msd deformable registration 0.48642522  
1005 1012  
    GDSC before registration: 0.25037813  
    GDSC after affine registration 0.4019997  
    GDSC after grad_msd deformable registration 0.45774922  
  
WEAK SEGMENTATIONS  
    Average GDSC before registration: (0.25532487, 0.03496879)
```


Average GDSC after affine registration: (0.3745071, 0.016430398)
Average GDSC after deformable registration: (0.47622523, 0.021808852)

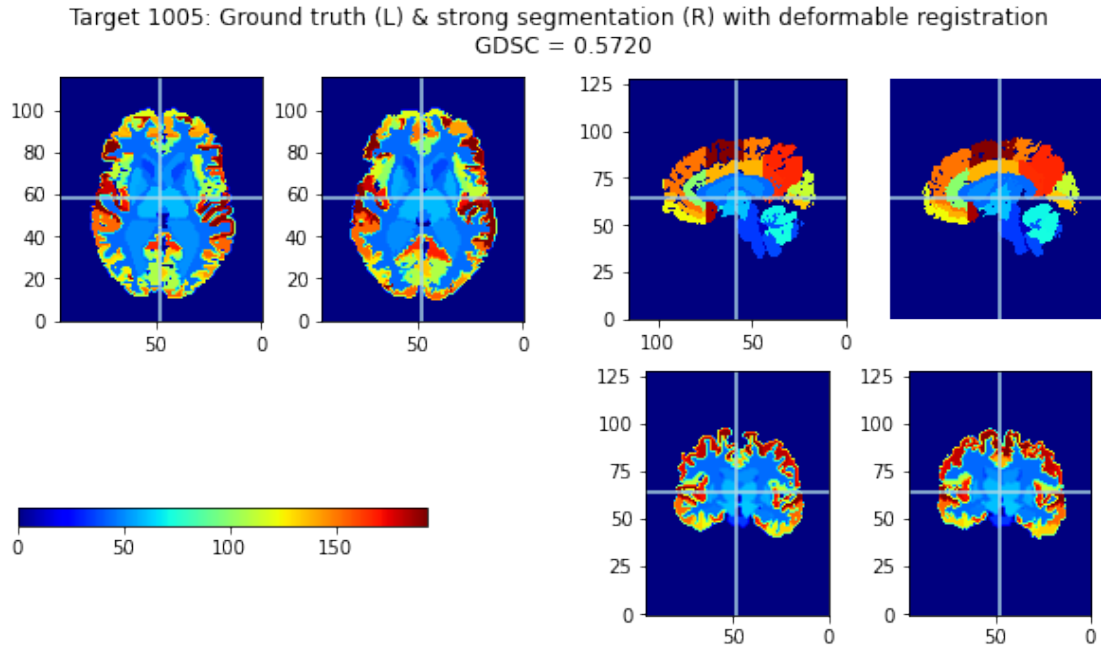
STRONG SEGMENTATIONS

GDSC with only affine registration: 0.45553008
GDSC with deformable registration: 0.5719938

```
[43]: my_view_multi([S_fix, strong_seg_aff], header=exp.hdr_fix, cmap='jet',
                    title=f"Target 1005: Ground truth (L) & strong segmentation (R)␣
                    ↳with only affine registration \nGDSC = {aff_gdsc:.4f}")

my_view_multi([S_fix, strong_seg_def], header=exp.hdr_fix, cmap='jet',
              title=f"Target 1005: Ground truth (L) & strong segmentation (R)␣
              ↳with deformable registration \nGDSC = {def_gdsc:.4f}")
```





4.4 Target Image 1018

```
[44]: strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, S_fix = get\_strong\_segmentation\(1018\)
```

```
1018 1000
GDSC before registration: 0.1224667
GDSC after affine registration 0.3658713
GDSC after grad_msd deformable registration 0.51392365
1018 1036
GDSC before registration: 0.1339297
GDSC after affine registration 0.36231488
GDSC after grad_msd deformable registration 0.46119422
1018 1001
GDSC before registration: 0.26013175
GDSC after affine registration 0.40272295
GDSC after grad_msd deformable registration 0.44714528
1018 1001
GDSC before registration: 0.26013175
GDSC after affine registration 0.40272295
GDSC after grad_msd deformable registration 0.44714528
1018 1036
GDSC before registration: 0.1339297
GDSC after affine registration 0.36231488
GDSC after grad_msd deformable registration 0.46119422
```

```

1018 1014
    GDSC before registration: 0.38094145
    GDSC after affine registration 0.41549993
    GDSC after grad_msd deformable registration 0.50124466
1018 1014
    GDSC before registration: 0.38094145
    GDSC after affine registration 0.41549993
    GDSC after grad_msd deformable registration 0.50124466
1018 1009
    GDSC before registration: 0.32738164
    GDSC after affine registration 0.38014004
    GDSC after grad_msd deformable registration 0.5111169
1018 1007
    GDSC before registration: 0.33075002
    GDSC after affine registration 0.4072966
    GDSC after grad_msd deformable registration 0.5107064
1018 1012
    GDSC before registration: 0.2794172
    GDSC after affine registration 0.42794567
    GDSC after grad_msd deformable registration 0.5113146

```

WEAK SEGMENTATIONS

```

Average GDSC before registration: (0.26100212, 0.09479947)
Average GDSC after affine registration: (0.39423293, 0.023230271)
Average GDSC after deformable registration: (0.48662296, 0.027152855)

```

STRONG SEGMENTATIONS

```

GDSC with only affine registration: 0.48106903
GDSC with deformable registration: 0.5740177

```

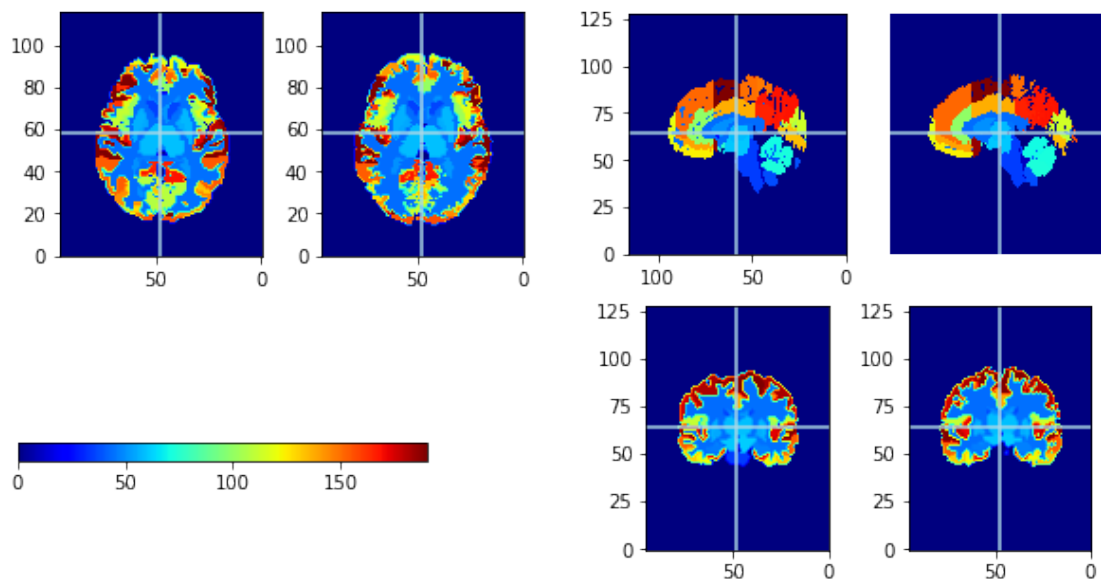
```

[45]: my_view_multi([S_fix, strong_seg_aff], header=exp.hdr_fix, cmap='jet',
                    title=f"Target 1018: Ground truth (L) & strong segmentation (R)␣
                    ↳with only affine registration \nGDSC = {aff_gdsc:.4f}")

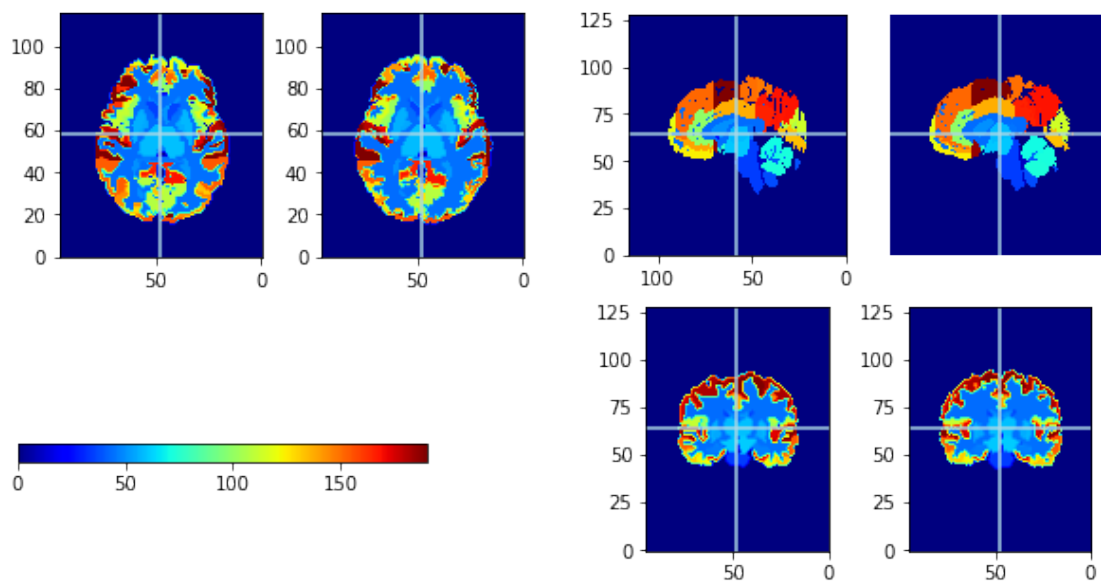
my_view_multi([S_fix, strong_seg_def], header=exp.hdr_fix, cmap='jet',
              title=f"Target 1018: Ground truth (L) & strong segmentation (R)␣
              ↳with deformable registration \nGDSC = {def_gdsc:.4f}")

```


Target 1018: Ground truth (L) & strong segmentation (R) with only affine registration
GDSC = 0.4811



Target 1018: Ground truth (L) & strong segmentation (R) with deformable registration
GDSC = 0.5740



4.5 Target Image 1019

```
[46]: strong_seg_aff, strong_seg_def, aff_gdsc, def_gdsc, S_fix =   
      ↪ get_strong_segmentation(1019)
```

```
1019 1000  
      GDSC before registration: 0.027663497  
      GDSC after affine registration 0.3542557  
      GDSC after grad_msd deformable registration 0.4780594  
1019 1036  
      GDSC before registration: 0.033684824  
      GDSC after affine registration 0.34974688  
      GDSC after grad_msd deformable registration 0.51479894  
1019 1001  
      GDSC before registration: 0.08404672  
      GDSC after affine registration 0.37625387  
      GDSC after grad_msd deformable registration 0.49218595  
1019 1001  
      GDSC before registration: 0.08404672  
      GDSC after affine registration 0.37625387  
      GDSC after grad_msd deformable registration 0.49218595  
1019 1036  
      GDSC before registration: 0.033684824  
      GDSC after affine registration 0.34974688  
      GDSC after grad_msd deformable registration 0.51479894  
1019 1014  
      GDSC before registration: 0.120412685  
      GDSC after affine registration 0.3365032  
      GDSC after grad_msd deformable registration 0.476658  
1019 1014  
      GDSC before registration: 0.120412685  
      GDSC after affine registration 0.3365032  
      GDSC after grad_msd deformable registration 0.476658  
1019 1009  
      GDSC before registration: 0.10763356  
      GDSC after affine registration 0.36100587  
      GDSC after grad_msd deformable registration 0.43466574  
1019 1007  
      GDSC before registration: 0.14144461  
      GDSC after affine registration 0.3597839  
      GDSC after grad_msd deformable registration 0.4779459  
1019 1012  
      GDSC before registration: 0.0719865  
      GDSC after affine registration 0.37483683  
      GDSC after grad_msd deformable registration 0.49617532
```

WEAK SEGMENTATIONS

```
Average GDSC before registration: (0.082501665, 0.038560484)
```


Average GDSC after affine registration: (0.35748905, 0.014261714)
Average GDSC after deformable registration: (0.48541322, 0.021847328)

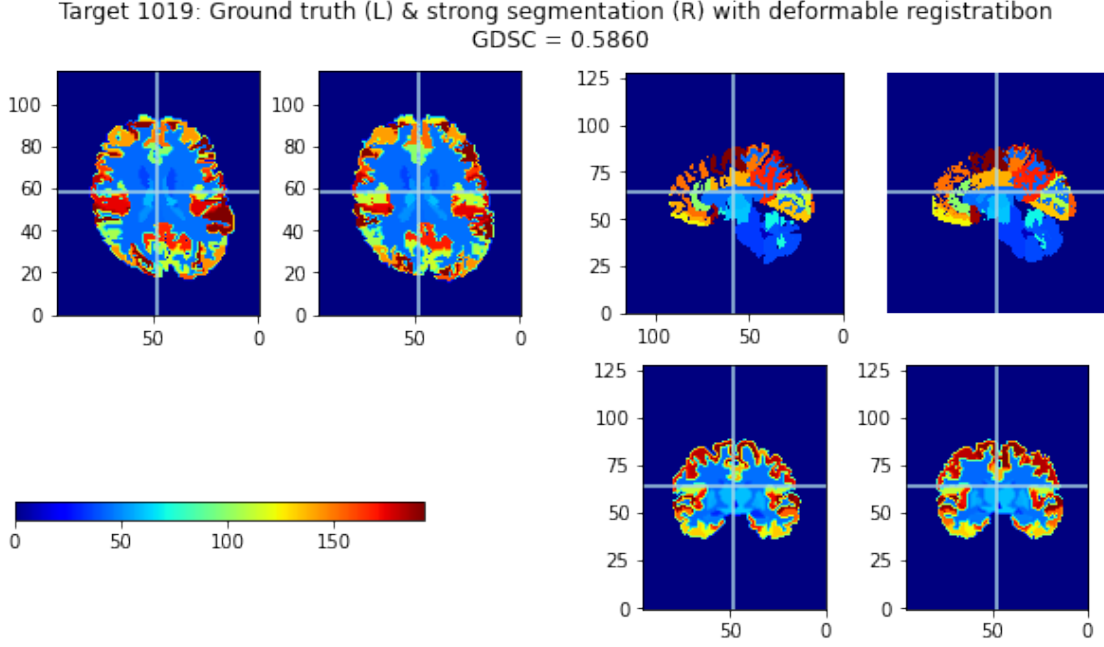
STRONG SEGMENTATIONS

GDSC with only affine registration: 0.42312393
GDSC with deformable registration: 0.5860492

```
[47]: my_view_multi([S_fix, strong_seg_aff], header=exp.hdr_fix, cmap='jet',
                    title=f"Target 1019: Ground truth (L) & strong segmentation (R)␣
                    ↳with only affine registration \nGDSC = {aff_gdsc:.4f}")

my_view_multi([S_fix, strong_seg_def], header=exp.hdr_fix, cmap='jet',
              title=f"Target 1019: Ground truth (L) & strong segmentation (R)␣
              ↳with deformable registratibon \nGDSC = {def_gdsc:.4f}")
```





4.6 Analysis

Target image	Avg GDSC weak affine	Avg GDSC weak deformable	GDSC strong affine	GDSC strong deformable
1003	0.3999 ± 0.0251	0.4221 ± 0.0602	0.4896	0.4916
1004	0.3960 ± 0.0223	0.4919 ± 0.0208	0.4855	0.5953
1005	0.3745 ± 0.0164	0.4762 ± 0.0218	0.4555	0.5720
1018	0.3942 ± 0.0232	0.4866 ± 0.0272	0.4811	0.5740
1019	0.3575 ± 0.0143	0.4854 ± 0.0218	0.4231	0.5860

Change from affine to deformable (strong)	Change from weak to strong (affine)	Change from weak to strong (deformable)
0.0020	0.0897	0.0695
0.1098	0.0895	0.1034
0.1165	0.0810	0.0958
0.0929	0.0869	0.0874
0.1629	0.0656	0.1006

For all images, we see an upward trend in the GDSC from affine to deformable registration. There is also gain in GDSC when we use the strong segmentation as compared to each of the weak segmentations.